Modeling with

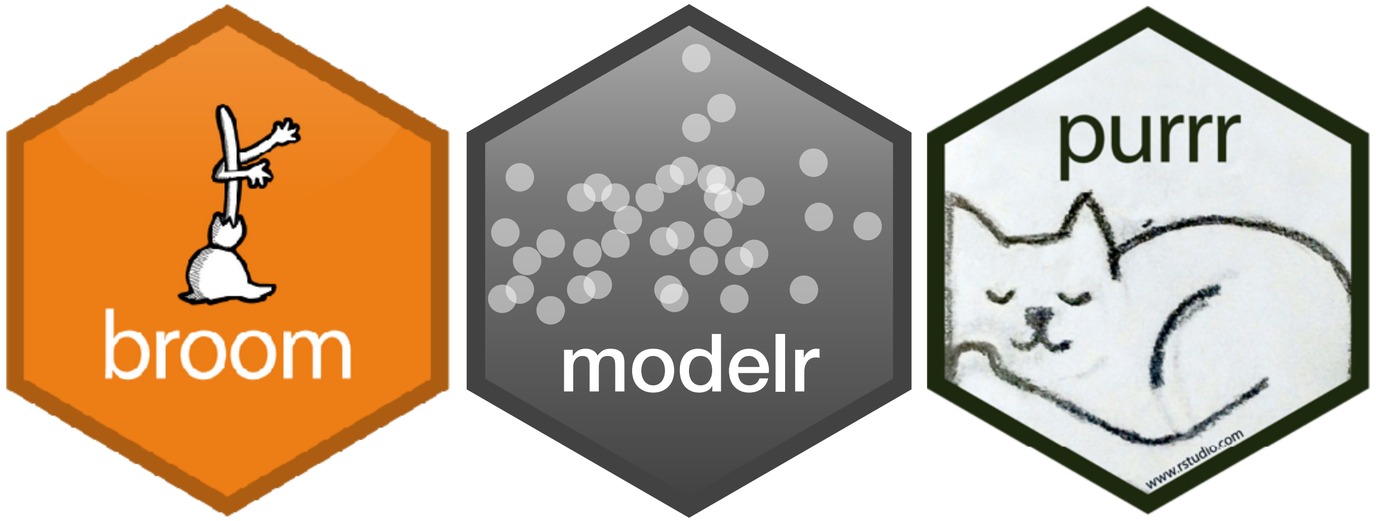

Open **08-List-Columns.Rmd**

gapminder

A subset of the data available at Hans Rosling's gapminder.org

```
# install.packages("gapminder")
library(gapminder)
```

gapminder

country <fctr>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>
Afghanistan	Asia	1952	28.80100	8425333	779.4453
Afghanistan	Asia	1957	30.33200	9240934	820.8530
Afghanistan	Asia	1962	31.99700	10267083	853.1007
Afghanistan	Asia	1967	34.02000	11537966	836.1971
Afghanistan	Asia	1972	36.08800	13079460	739.9811
Afghanistan	Asia	1977	38.43800	14880372	786.1134
Afghanistan	Asia	1982	39.85400	12881816	978.0114
Afghanistan	Asia	1987	40.82200	13867957	852.3959
Afghanistan	Asia	1992	41.67400	16317921	649.3414
Afghanistan	Asia	1997	41.76300	22227415	635.3414

1–10 of 1,704 rows Previous 1 2 3 4 5 6 ... 100 Next

Your Turn 1

How has life expectancy changed since 1952?

Make a line plot of **lifeExp** vs. **year** grouped by **country**. Set alpha to 0.2, to see the results better.

```
gapminder %>%
  ggplot(mapping = aes(x = year, y = lifeExp, group = country)) +
    geom_line(alpha = 0.2)
```

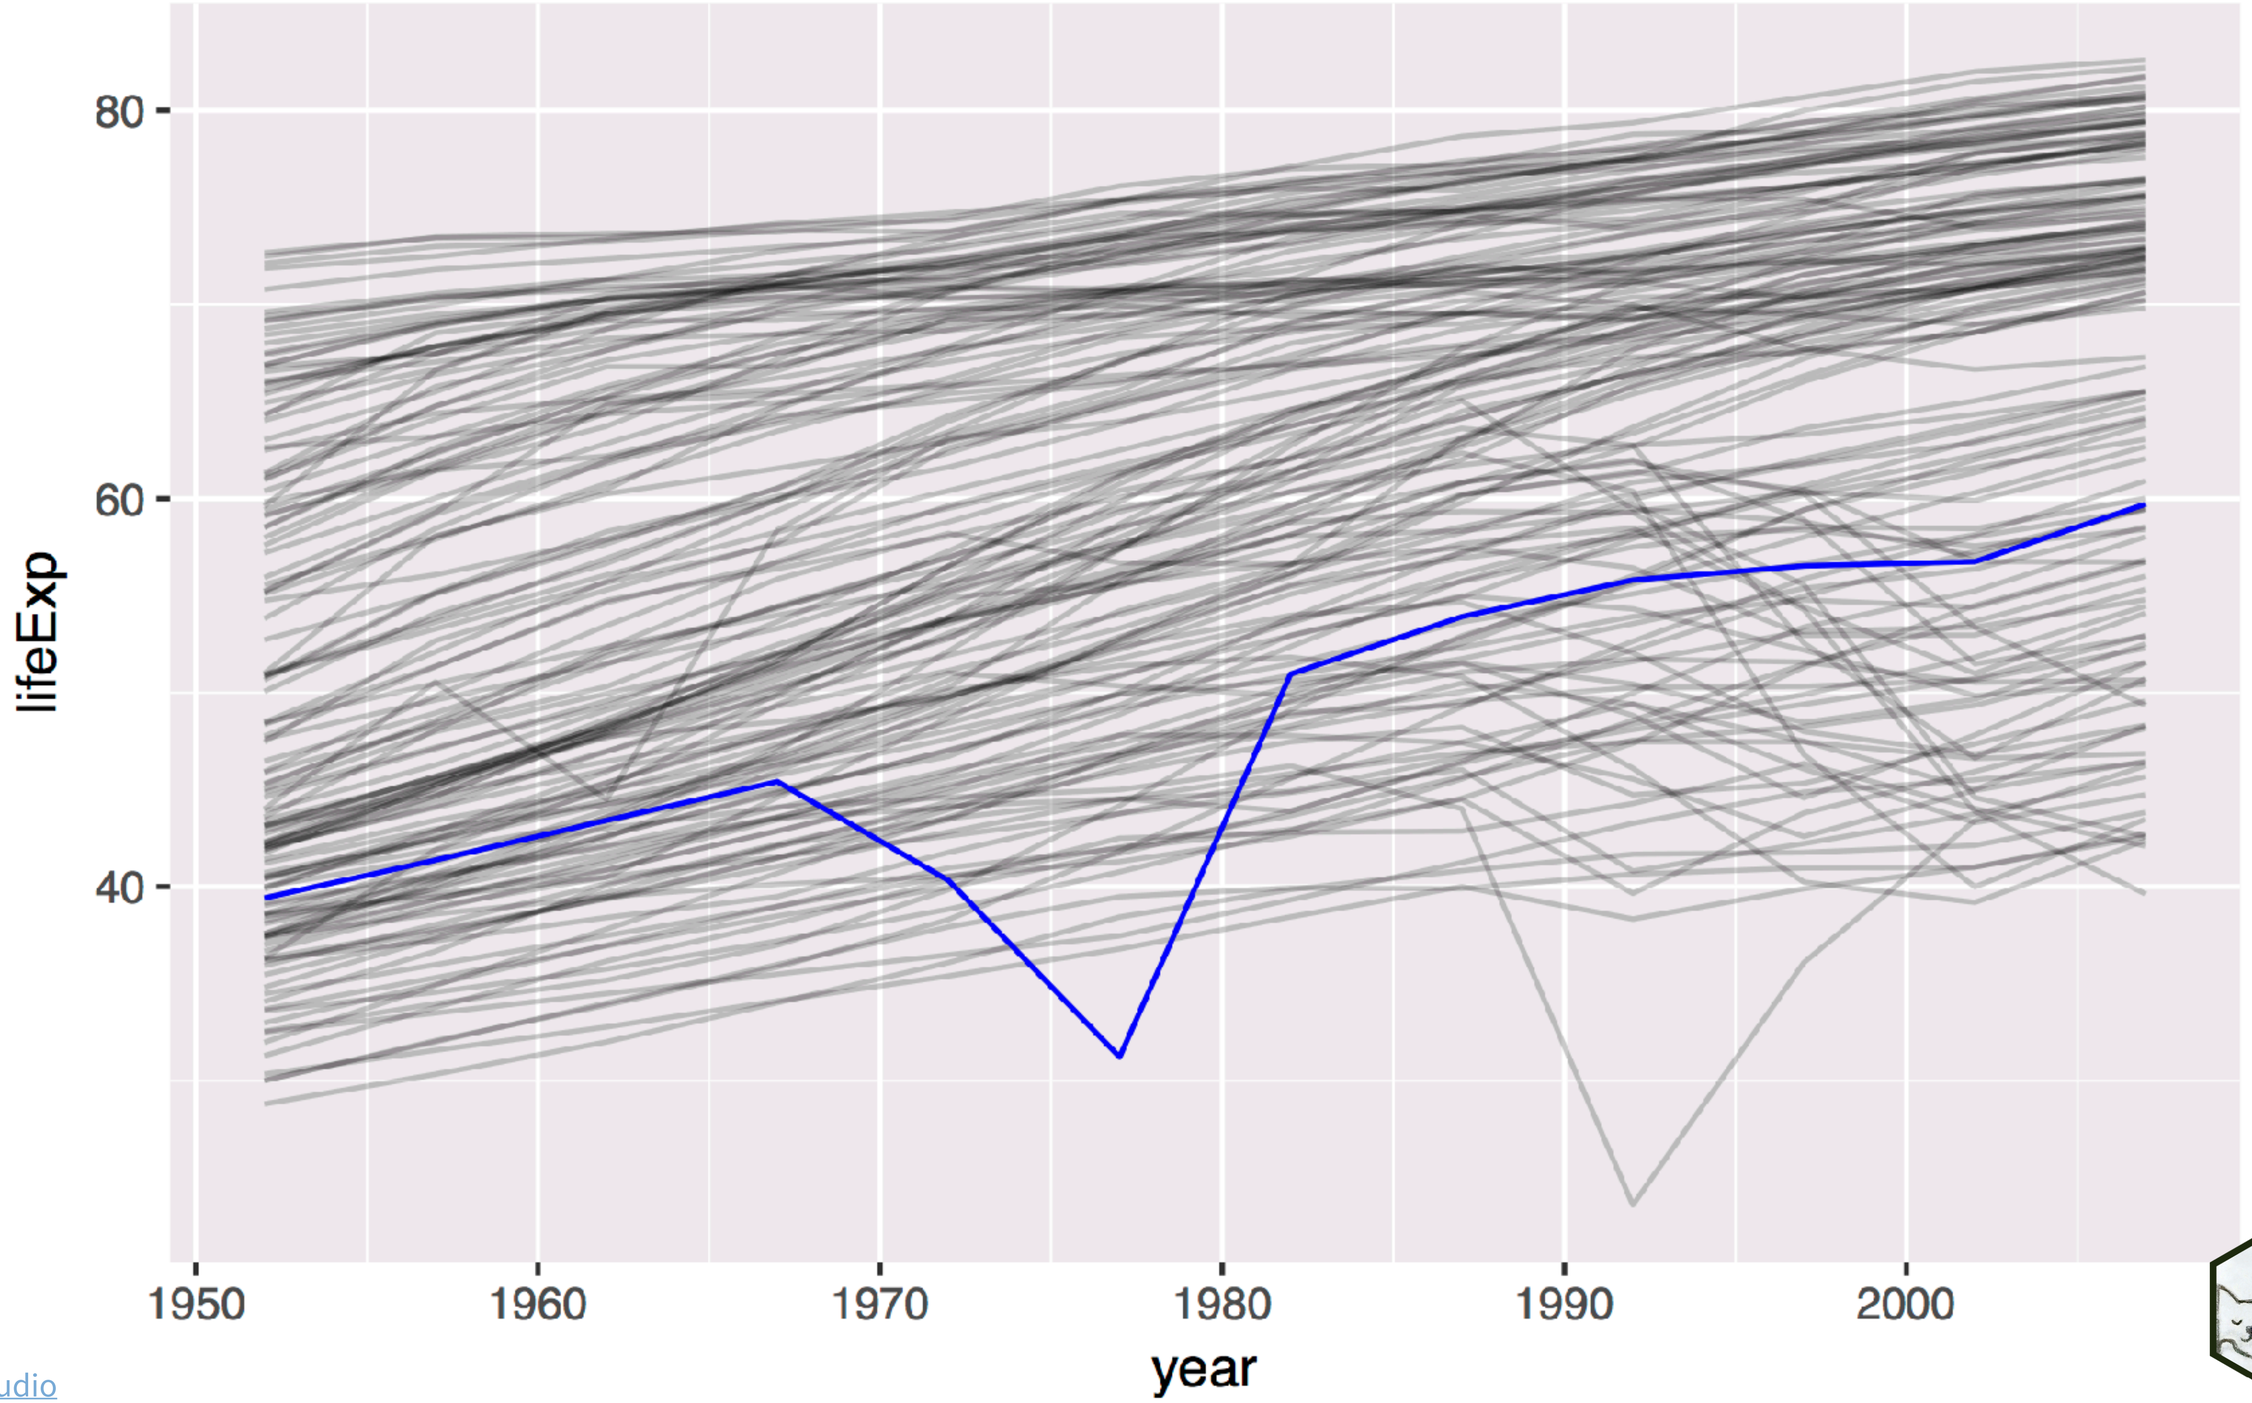

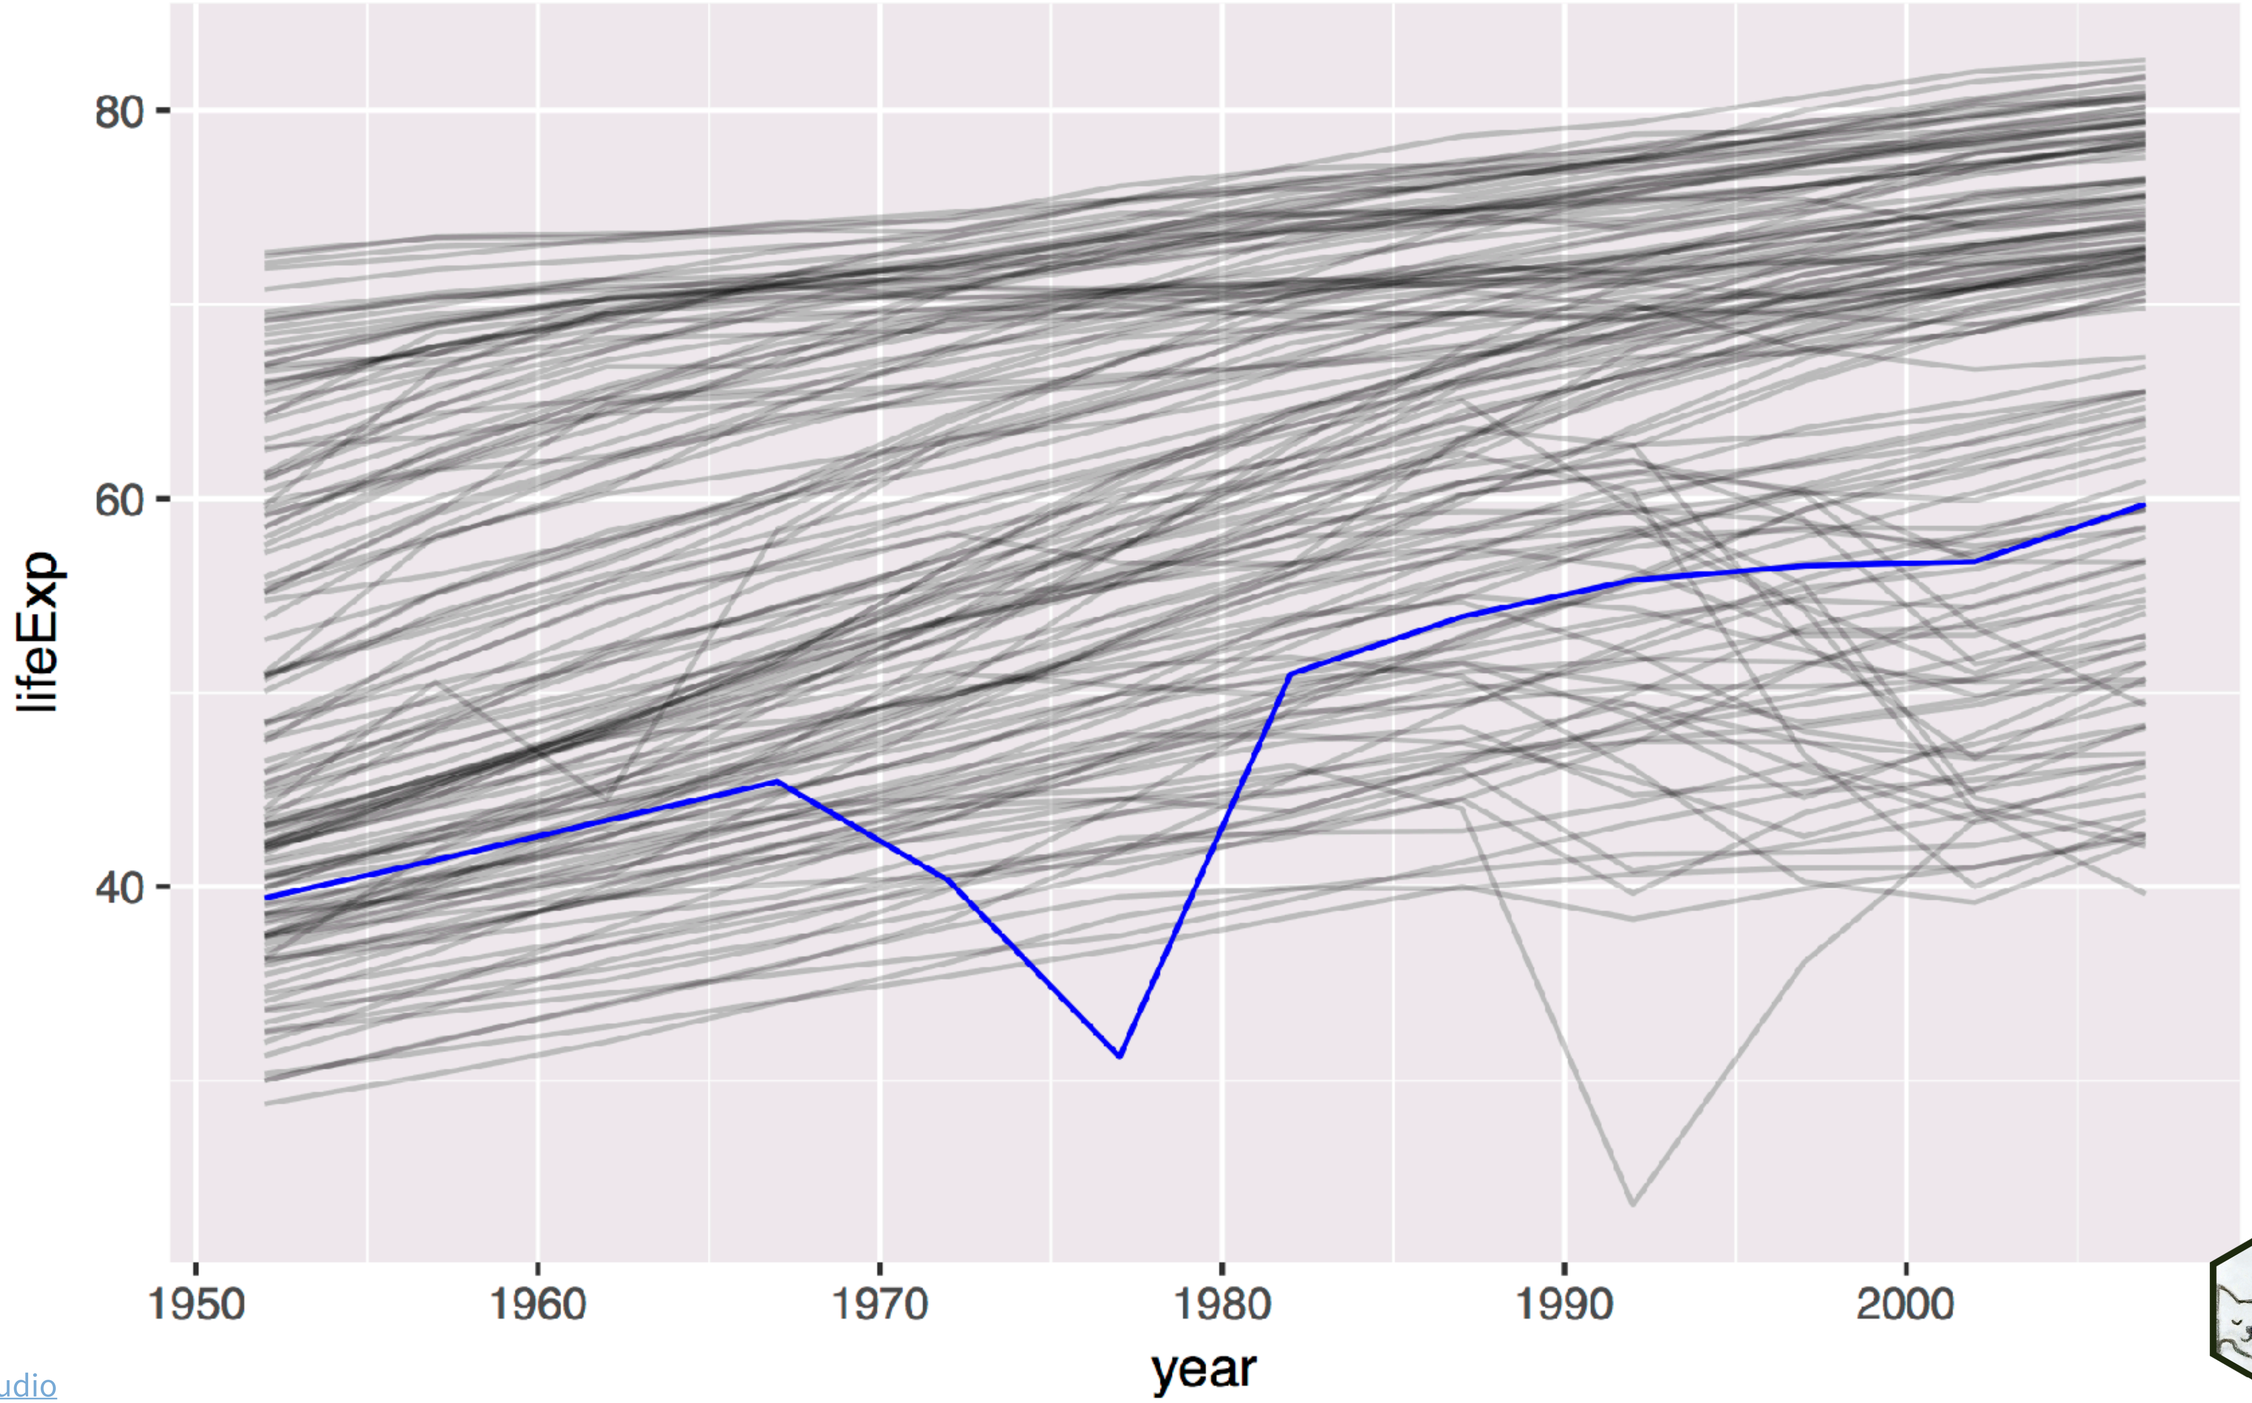

Idea 1

To quantify "linearity," fit a linear model, compare **r-squared**.

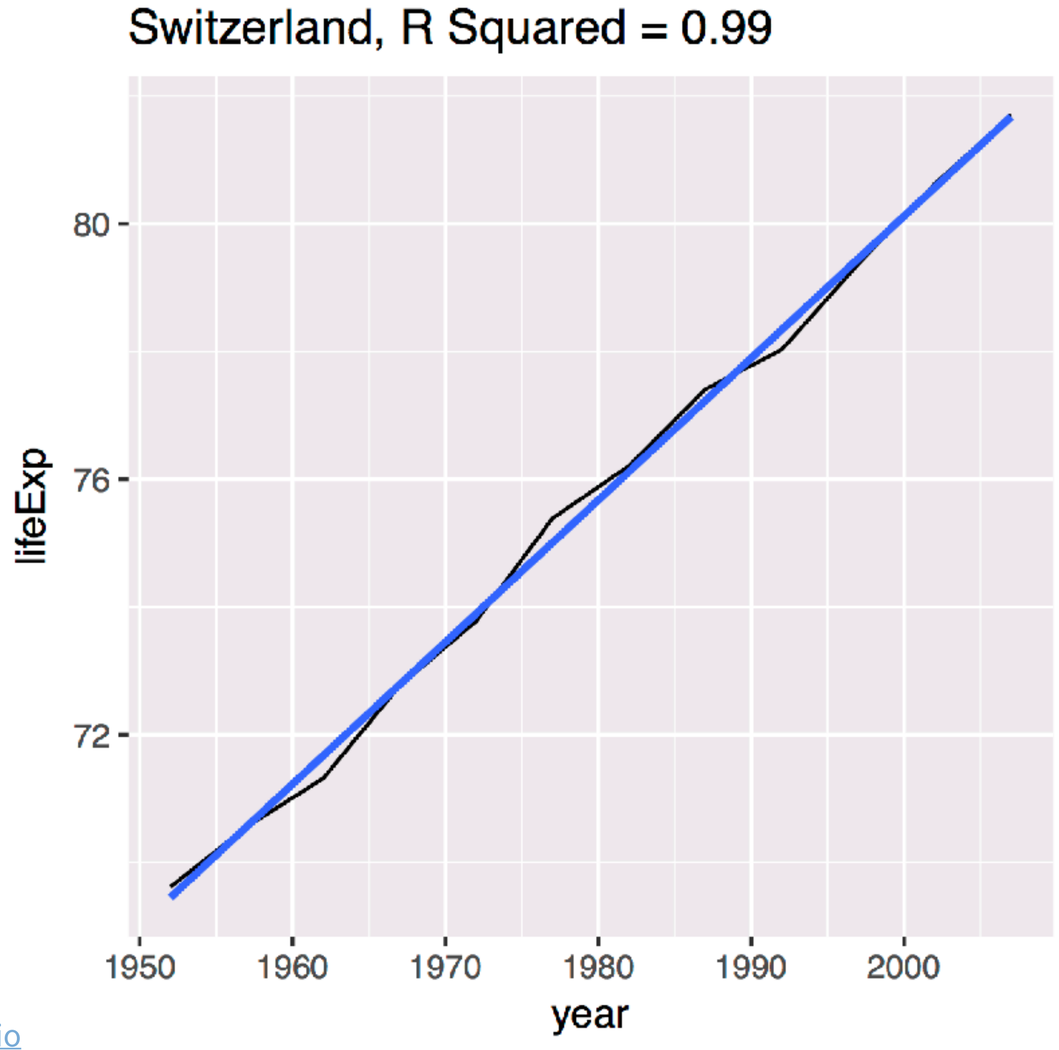

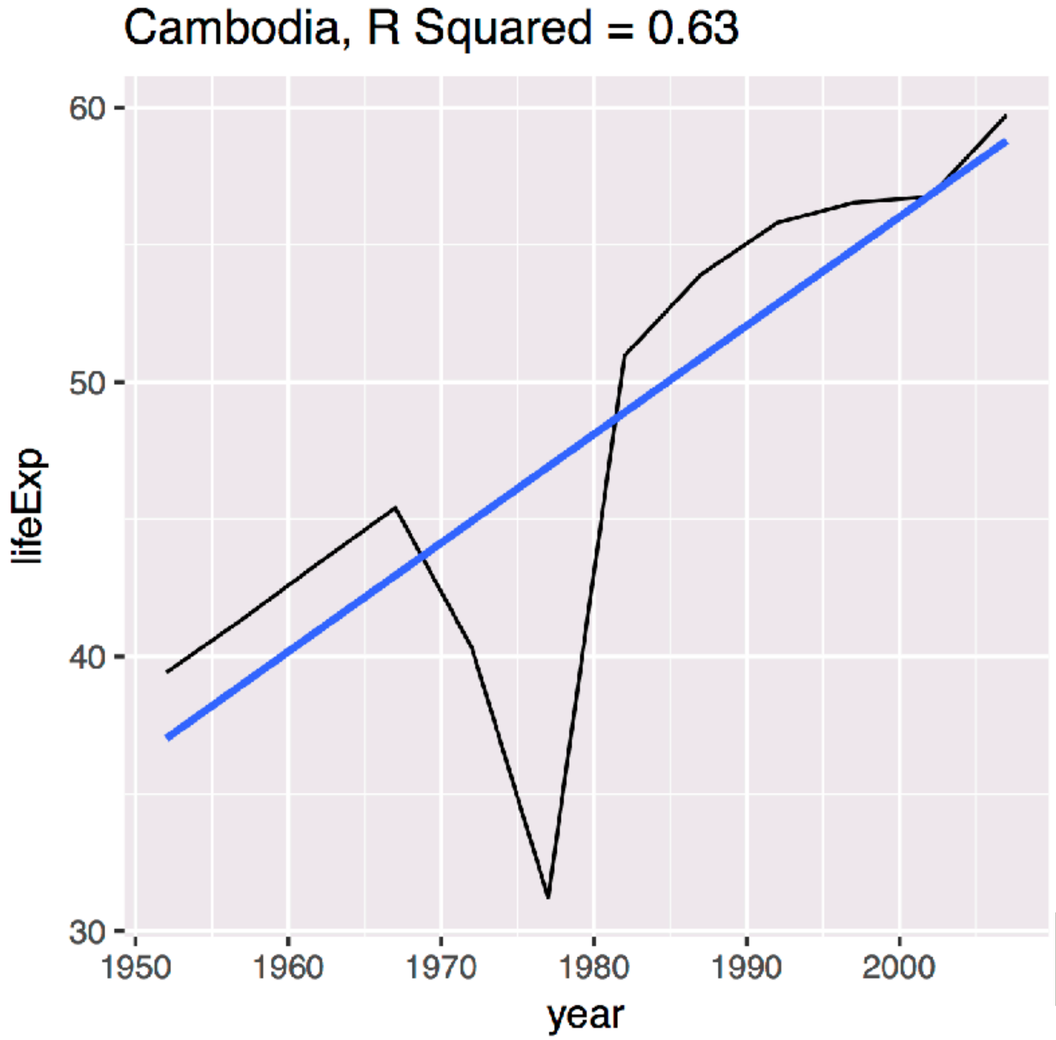

Idea 2

Use residuals to study deviations from the trend to study country specific patterns.

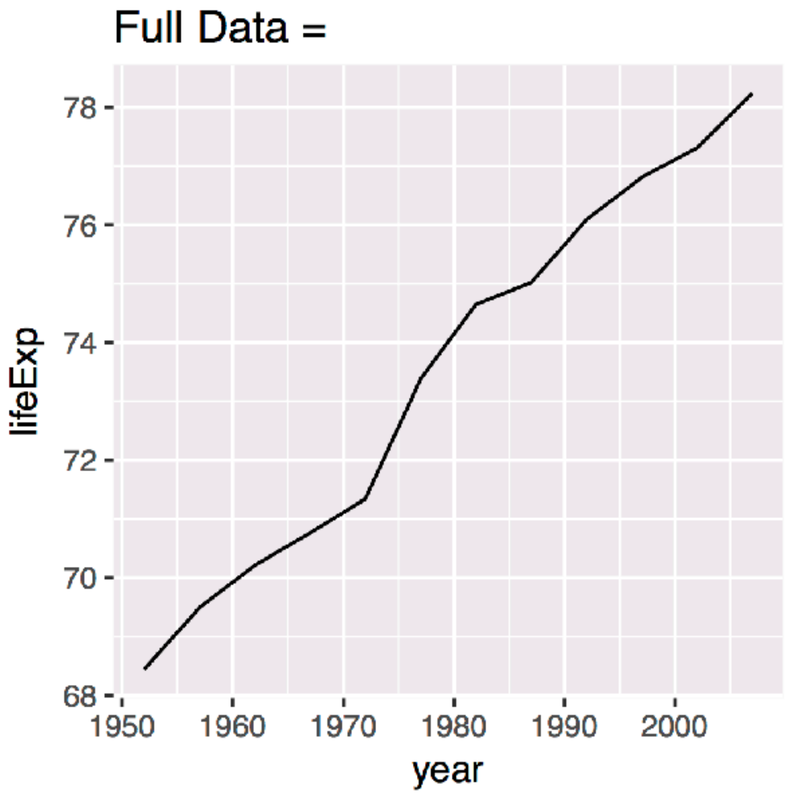

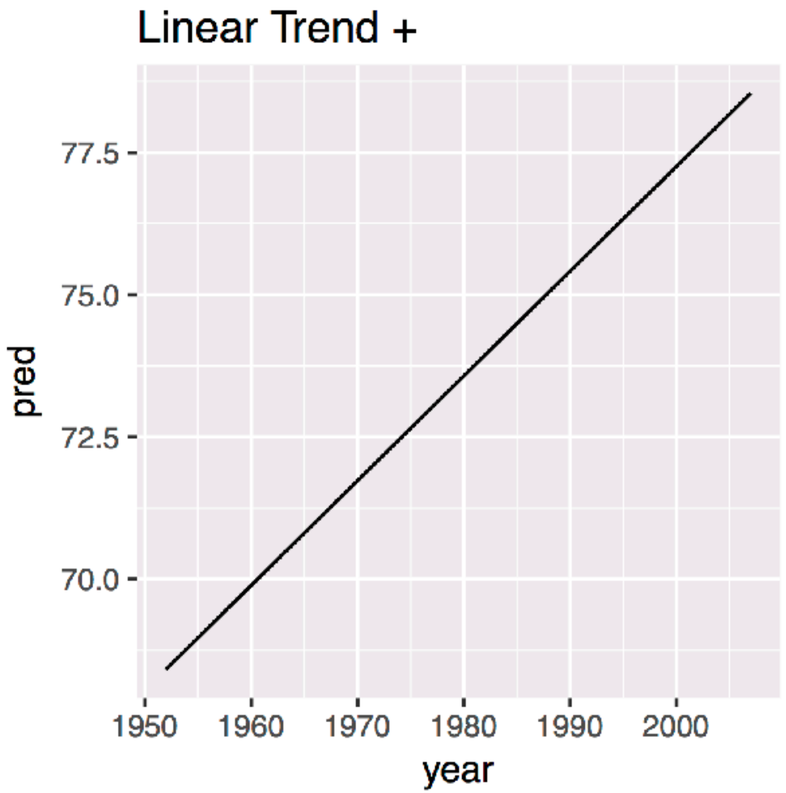

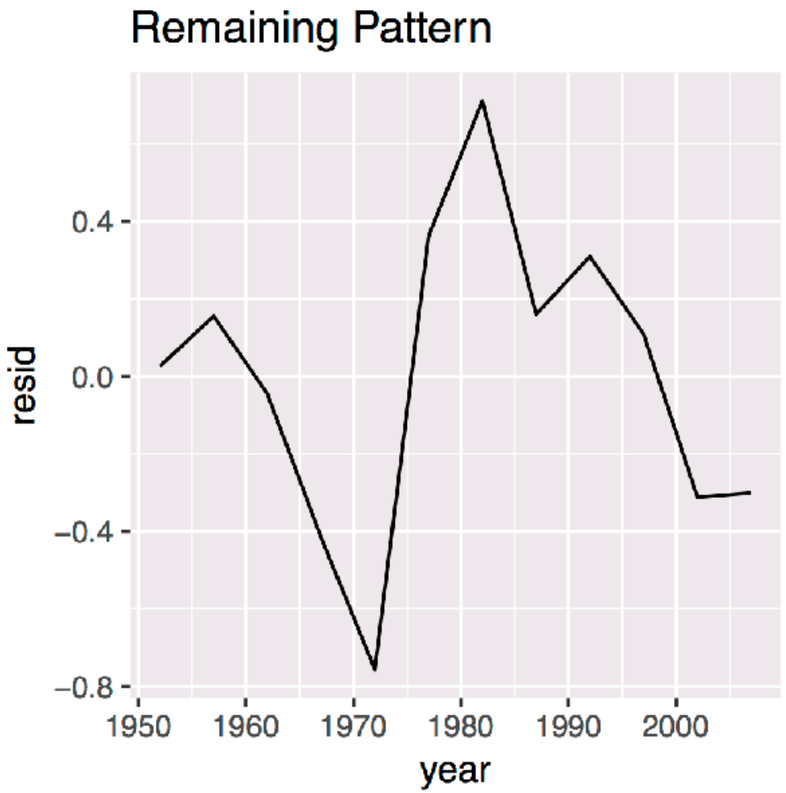

Idea 2

Use residuals to study deviations from the trend to study country specific patterns.

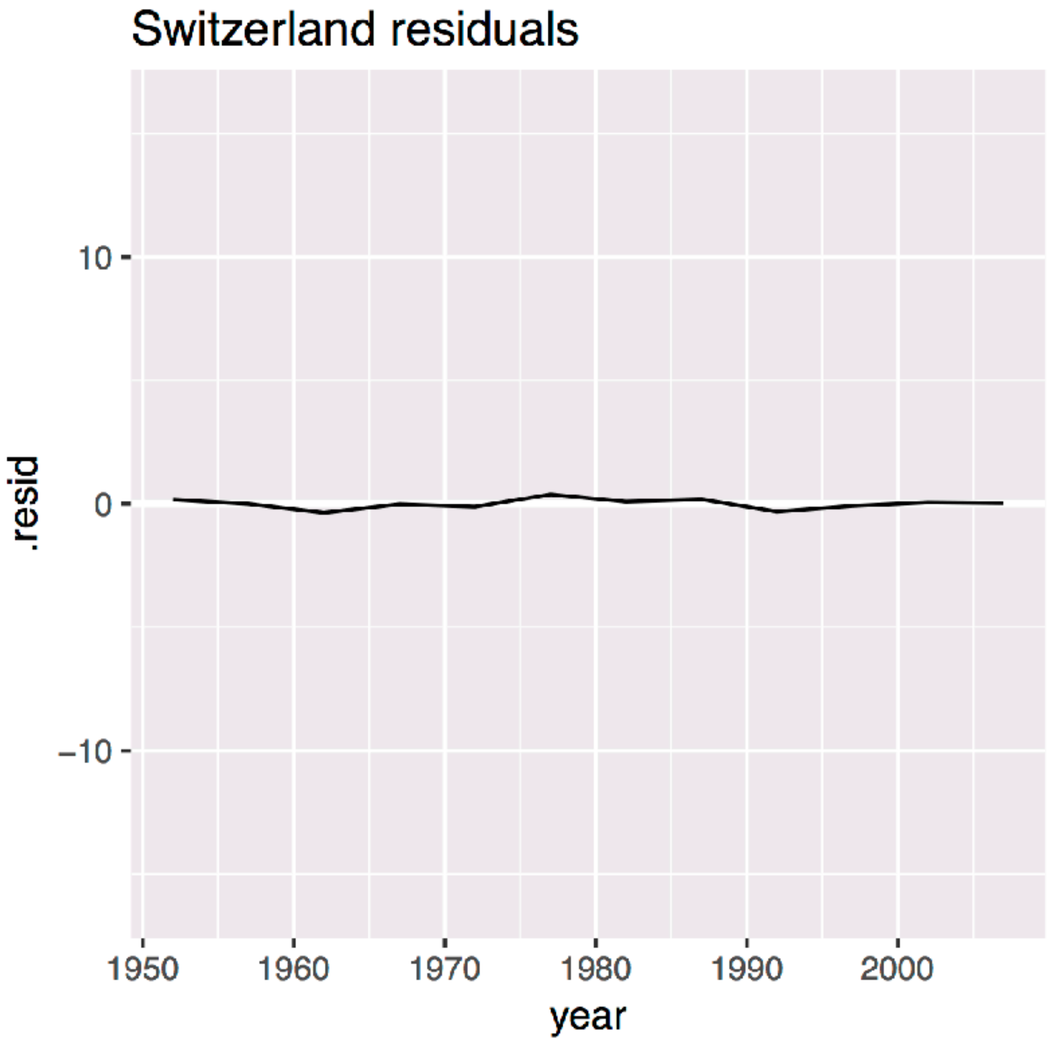

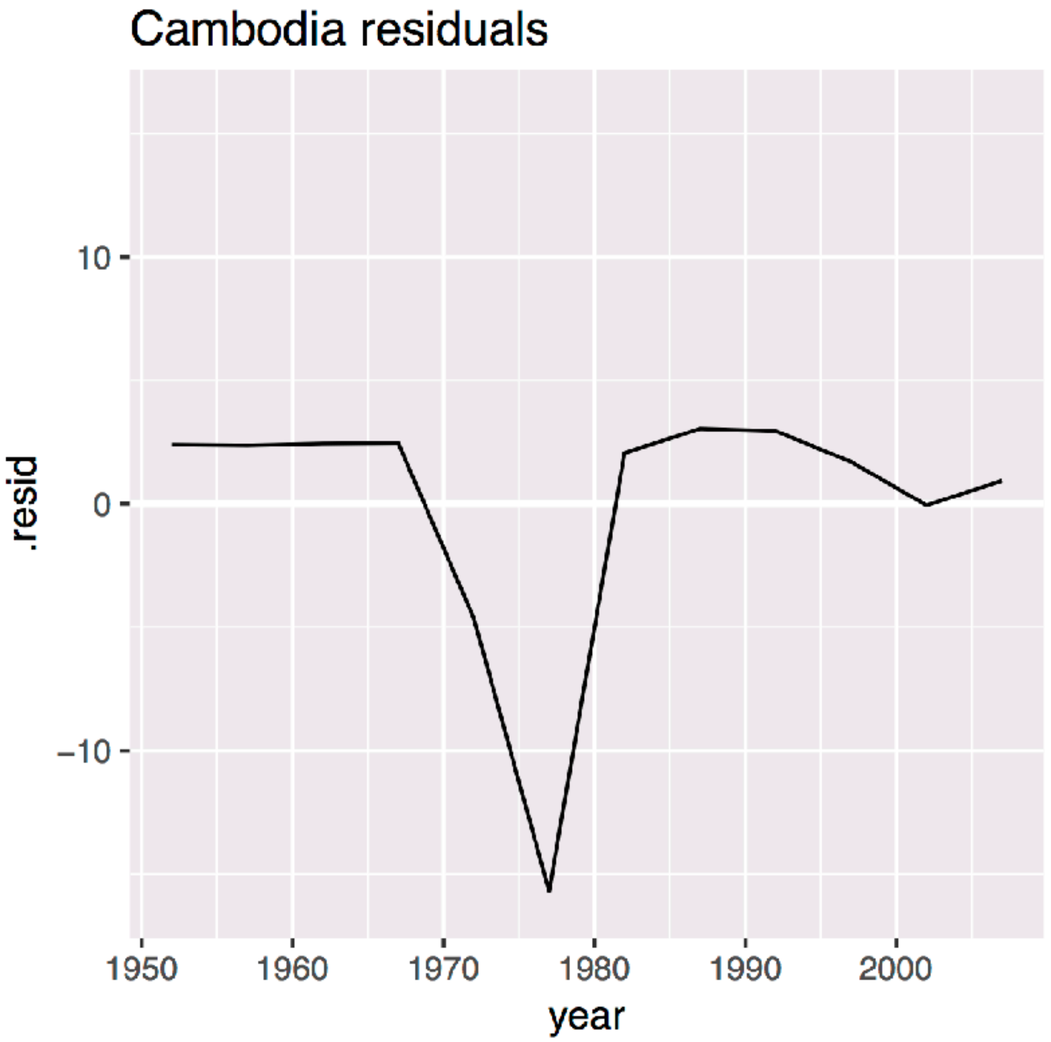

Goal

Fit model, compute r.squared, collect residuals ***for every country.***

1. **dplyr** grouping toolkit
2. **purrr** toolkit and list columns

dplyr do()

group_by() + do()

Run an expression on each group. The expression should return a data frame.

```
do(data, expression)
```

A grouped data frame

An R expression that returns a data frame

Use a "." in the expression to pass input (like a pipe)

group_by() + do()

Run an expression on each group. The expression should return a data frame.

```
gapminder %>%
  group_by(country, continent) %>%
  do()
```

A grouped data frame

group_by() + do()

Run an expression on each group. The expression should return a data frame.

```
gapminder %>%
  group_by(country, continent) %>%
  do(lm(lifeExp ~ year, data = .))
```

A grouped data frame

An R expression that uses a "." to pass input

but DOES NOT return a data frame

group_by() + do()

Run an expression on each group. The expression should return a data frame.

```
gapminder %>%
  group_by(country, continent) %>%
  do(glance(lm(lifeExp ~ year, data = .)))
```

A grouped data frame

An R expression that uses a "." to pass input

…and returns a data frame

group_by() + do()

Run an expression on each group. The expression should return a data frame.

```
gapminder %>%
  group_by(country, continent) %>%
  do(lm(lifeExp ~ year, data = .) %>% glance())
```

A grouped data frame

An R expression that uses a "." to pass input

…and returns a data frame

group_by() + do()

```
gapminder %>%
  group_by(country, continent) %>%
  do(lm(lifeExp ~ year, data = .) %>% glance())
```

country <fctr>	continent <fctr>	r.squared <dbl>	adj.r.squared <dbl>	sigma <dbl>	statistic <dbl>
Afghanistan	Asia	0.94771226	0.942483483	1.2227880	181.2494098
Albania	Europe	0.91057777	0.901635545	1.9830615	101.8290138
Algeria	Africa	0.98511721	0.983628932	1.3230064	661.9170864
Angola	Africa	0.88781463	0.876596093	1.4070091	79.1381823
Argentina	Americas	0.99556810	0.995124905	0.2923072	2246.3663487
Australia	Oceania	0.97964774	0.977612511	0.6206086	481.3458627
Austria	Europe	0.99213401	0.991347414	0.4074094	1261.2962902
Bahrain	Asia	0.96673981	0.963413791	1.6395865	290.6597394
Bangladesh	Asia	0.98936087	0.988296956	0.9766908	929.9263688
Belgium	Europe	0.99454056	0.993994612	0.2929025	1821.6883955

1–10 of 142 rows | 1–6 of 13 columns

Previous 1 2 3 4 5 6 ... 15 Next

group_by() + do()

```
gapminder %>%
  group_by(country, continent) %>%
  do(lm(lifeExp ~ year, data = .) %>% glance())
```

country <fctr>	continent <fctr>	r.squared <dbl>	adj.r.squared <dbl>	sigma <dbl>	statistic <dbl>
Afghanistan	Asia	0.94771226	0.942483483	1.2227880	181.2494098
Albania	Europe	0.91057777	0.901635545	1.9830615	101.8290138
Algeria	Africa	0.98511721	0.983628932	1.3230064	661.9170864
Angola	Africa	0.88781463	0.876596093	1.4070091	79.1381823
Argentina	Americas	0.99556810	0.995124905	0.2923072	2246.3663487
Australia	Oceania	0.97964774	0.977612511	0.6206086	481.3458627
Austria	Europe	0.99213401	0.991347414	0.4074094	1261.2962902
Bahrain	Asia	0.96673981	0.963413791	1.6395865	290.6597394
Bangladesh	Asia	0.98936087	0.988296956	0.9766908	929.9263688
Belgium	Europe	0.99454056	0.993994612	0.2929025	1821.6883955

1–10 of 142 rows | 1–6 of 13 columns

Previous 1 2 3 4 5 6 ... 15 Next

group_by() + do()

```
gapminder %>%
  group_by(country, continent) %>%
  do(lm(lifeExp ~ year, data = .) %>% tidy())
```

country <fctr>	continent <fctr>	term <chr>	estimate <dbl>	std.error <dbl>
Afghanistan	Asia	(Intercept)	-5.075343e+02	4.048416e+01
Afghanistan	Asia	year	2.753287e-01	2.045093e-02
Albania	Europe	(Intercept)	-5.940725e+02	6.565536e+01
Albania	Europe	year	3.346832e-01	3.316639e-02
Algeria	Africa	(Intercept)	-1.067859e+03	4.380220e+01
Algeria	Africa	year	5.692797e-01	2.212707e-02
Angola	Africa	(Intercept)	-3.765048e+02	4.658337e+01
Angola	Africa	year	2.093399e-01	2.353200e-02
Argentina	Americas	(Intercept)	-3.896063e+02	9.677730e+00
Argentina	Americas	year	2.317084e-01	4.888791e-03

1–10 of 284 rows | 1–5 of 7 columns

Previous 1 2 3 4 5 6 ... 29 Next

group_by() + do()

```
gapminder %>%
  group_by(country, continent) %>%
  do(lm(lifeExp ~ year, data = .) %>% augment())
```

country <fctr>	continent <fctr>	lifeExp <dbl>	year <int>	.fitted <dbl>	.se.fit <dbl>	.resid <dbl>	.hat <dbl>
Afghanistan	Asia	28.80100	1952	29.90729	0.66399954	-1.106295e+00	0.29487179
Afghanistan	Asia	30.33200	1957	31.28394	0.57994415	-9.519382e-01	0.22494172
Afghanistan	Asia	31.99700	1962	32.66058	0.50267990	-6.635816e-01	0.16899767
Afghanistan	Asia	34.02000	1967	34.03722	0.43583366	-1.722494e-02	0.12703963
Afghanistan	Asia	36.08800	1972	35.41387	0.38487259	6.741317e-01	0.09906760
Afghanistan	Asia	38.43800	1977	36.79051	0.35667194	1.647488e+00	0.08508159
Afghanistan	Asia	39.85400	1982	38.16716	0.35667194	1.686845e+00	0.08508159
Afghanistan	Asia	40.82200	1987	39.54380	0.38487259	1.278202e+00	0.09906760
Afghanistan	Asia	41.67400	1992	40.92044	0.43583366	7.535583e-01	0.12703963
Afghanistan	Asia	41.76300	1997	42.29709	0.50267990	-5.340851e-01	0.16899767

1–10 of 1,704 rows | 1–8 of 11 columns

Previous 1 2 3 4 5 6 ... 100 Next

Your Turn 2

Group the data by **country** and **continent** then fit a model and collect the residuals *for each country*.

Plot the **residuals** vs **year** as a line graph, grouped by **country**, with alpha = 0.2.

Add the following to your plot to facet by **continent**:

+ facet_wrap(~continent)

```
gapminder %>%
  group_by(country, continent) %>%
  do(lm(lifeExp ~ year, data = .) %>%  augment()) %>%
  ggplot() +
    geom_line(aes(year, .resid, group = country), alpha = 0.2) +
    facet_wrap(~continent)
```

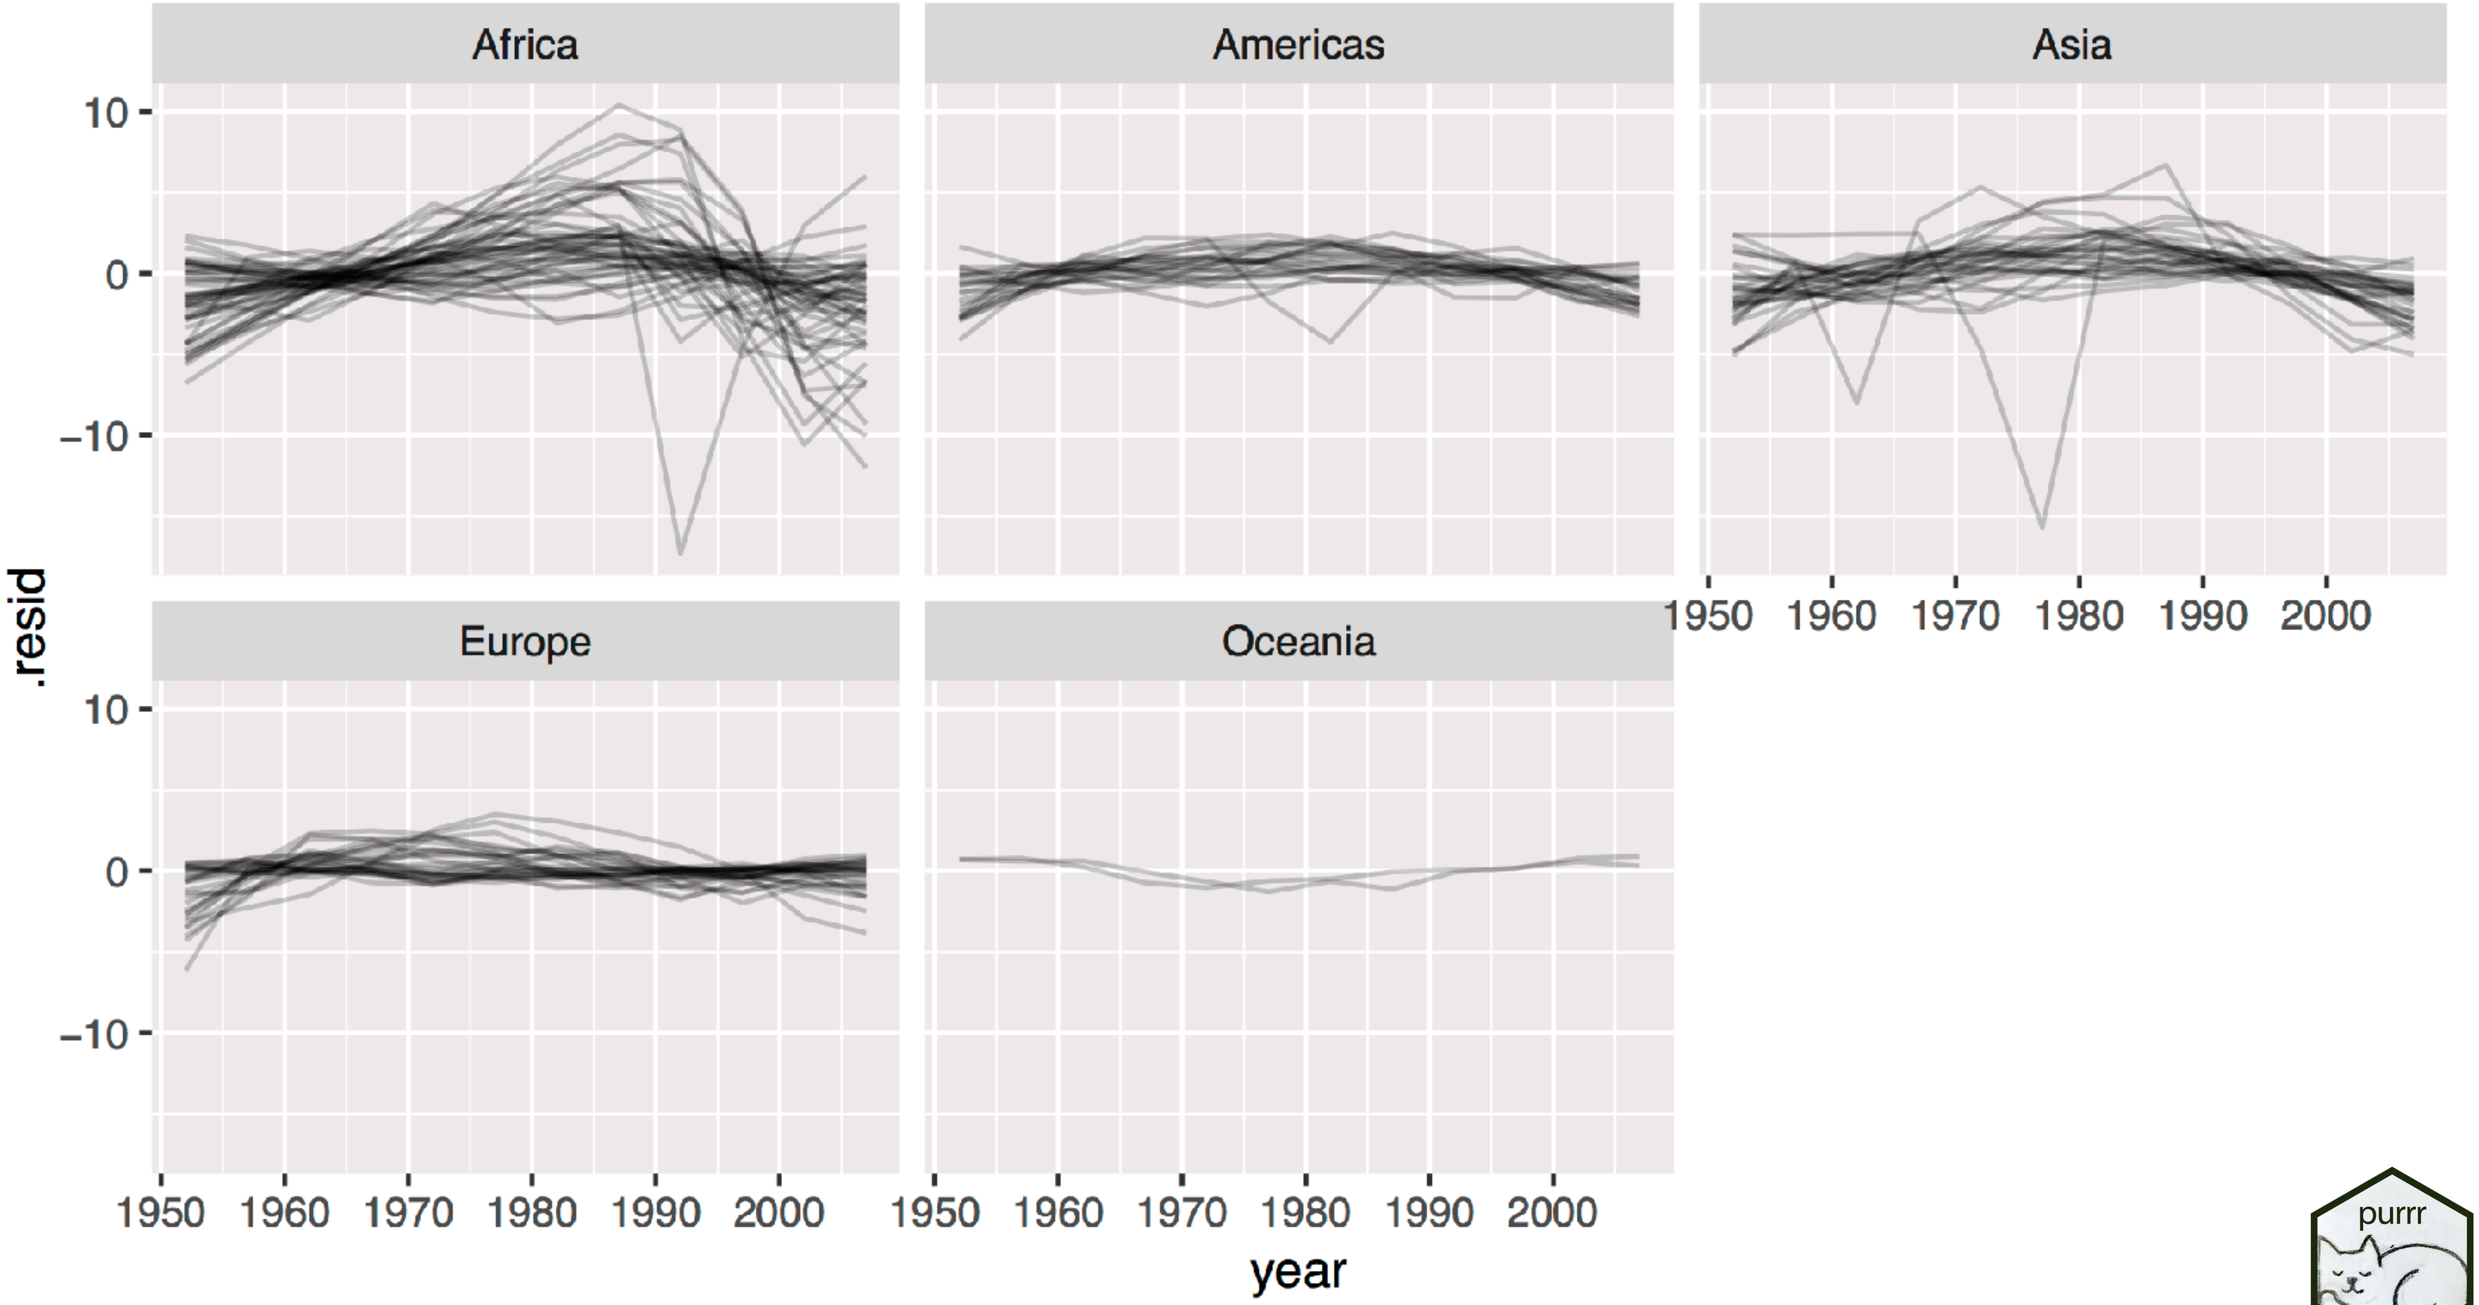

Quiz

Which broom function can we use to find the **r squared** for each model?

Bad fits

Which countries had a non-linear progress?

```
gapminder %>%
  group_by(country, continent) %>%
  do(lm(lifeExp ~ year, data = .) %>% glance()) %>%
  filter(r.squared < 0.25)
```

country <fctr>	continent <fctr>	r.squared <dbl>	adj.r.squared <dbl>	sigma <dbl>	statistic <dbl>
Botswana	Africa	0.03402340	-0.062574259	6.112177	0.3522177
Lesotho	Africa	0.08485635	-0.006658011	5.933934	0.9272463
Rwanda	Africa	0.01715964	-0.081124401	6.558269	0.1745923
Swaziland	Africa	0.06821087	-0.024968046	6.644091	0.7320419
Zambia	Africa	0.05983644	-0.034179918	4.528713	0.6364471
Zimbabwe	Africa	0.05623196	-0.038144842	7.205431	0.5958240

Your Turn 3

Complete the code to filter the dataset that you made in Your Turn 2 against **bad_fits**.

Use the result to plot a line graph of **year** vs. .**resid** colored by **country** for each country that had an r-squared < 0.25.

```
bad_fits <- gapminder %>%
  group_by(country) %>%
  do(lm(lifeExp ~ year, data = .) %>%  glance()) %>%
  filter(r.squared < 0.25)

residuals <- gapminder %>%
  group_by(country) %>%
  do(lm(lifeExp ~ year, data = .) %>%  augment())

residuals %>%
  semi_join(bad_fits) %>%
  ggplot() +
    geom_line(aes(year, .resid, color = country))
```

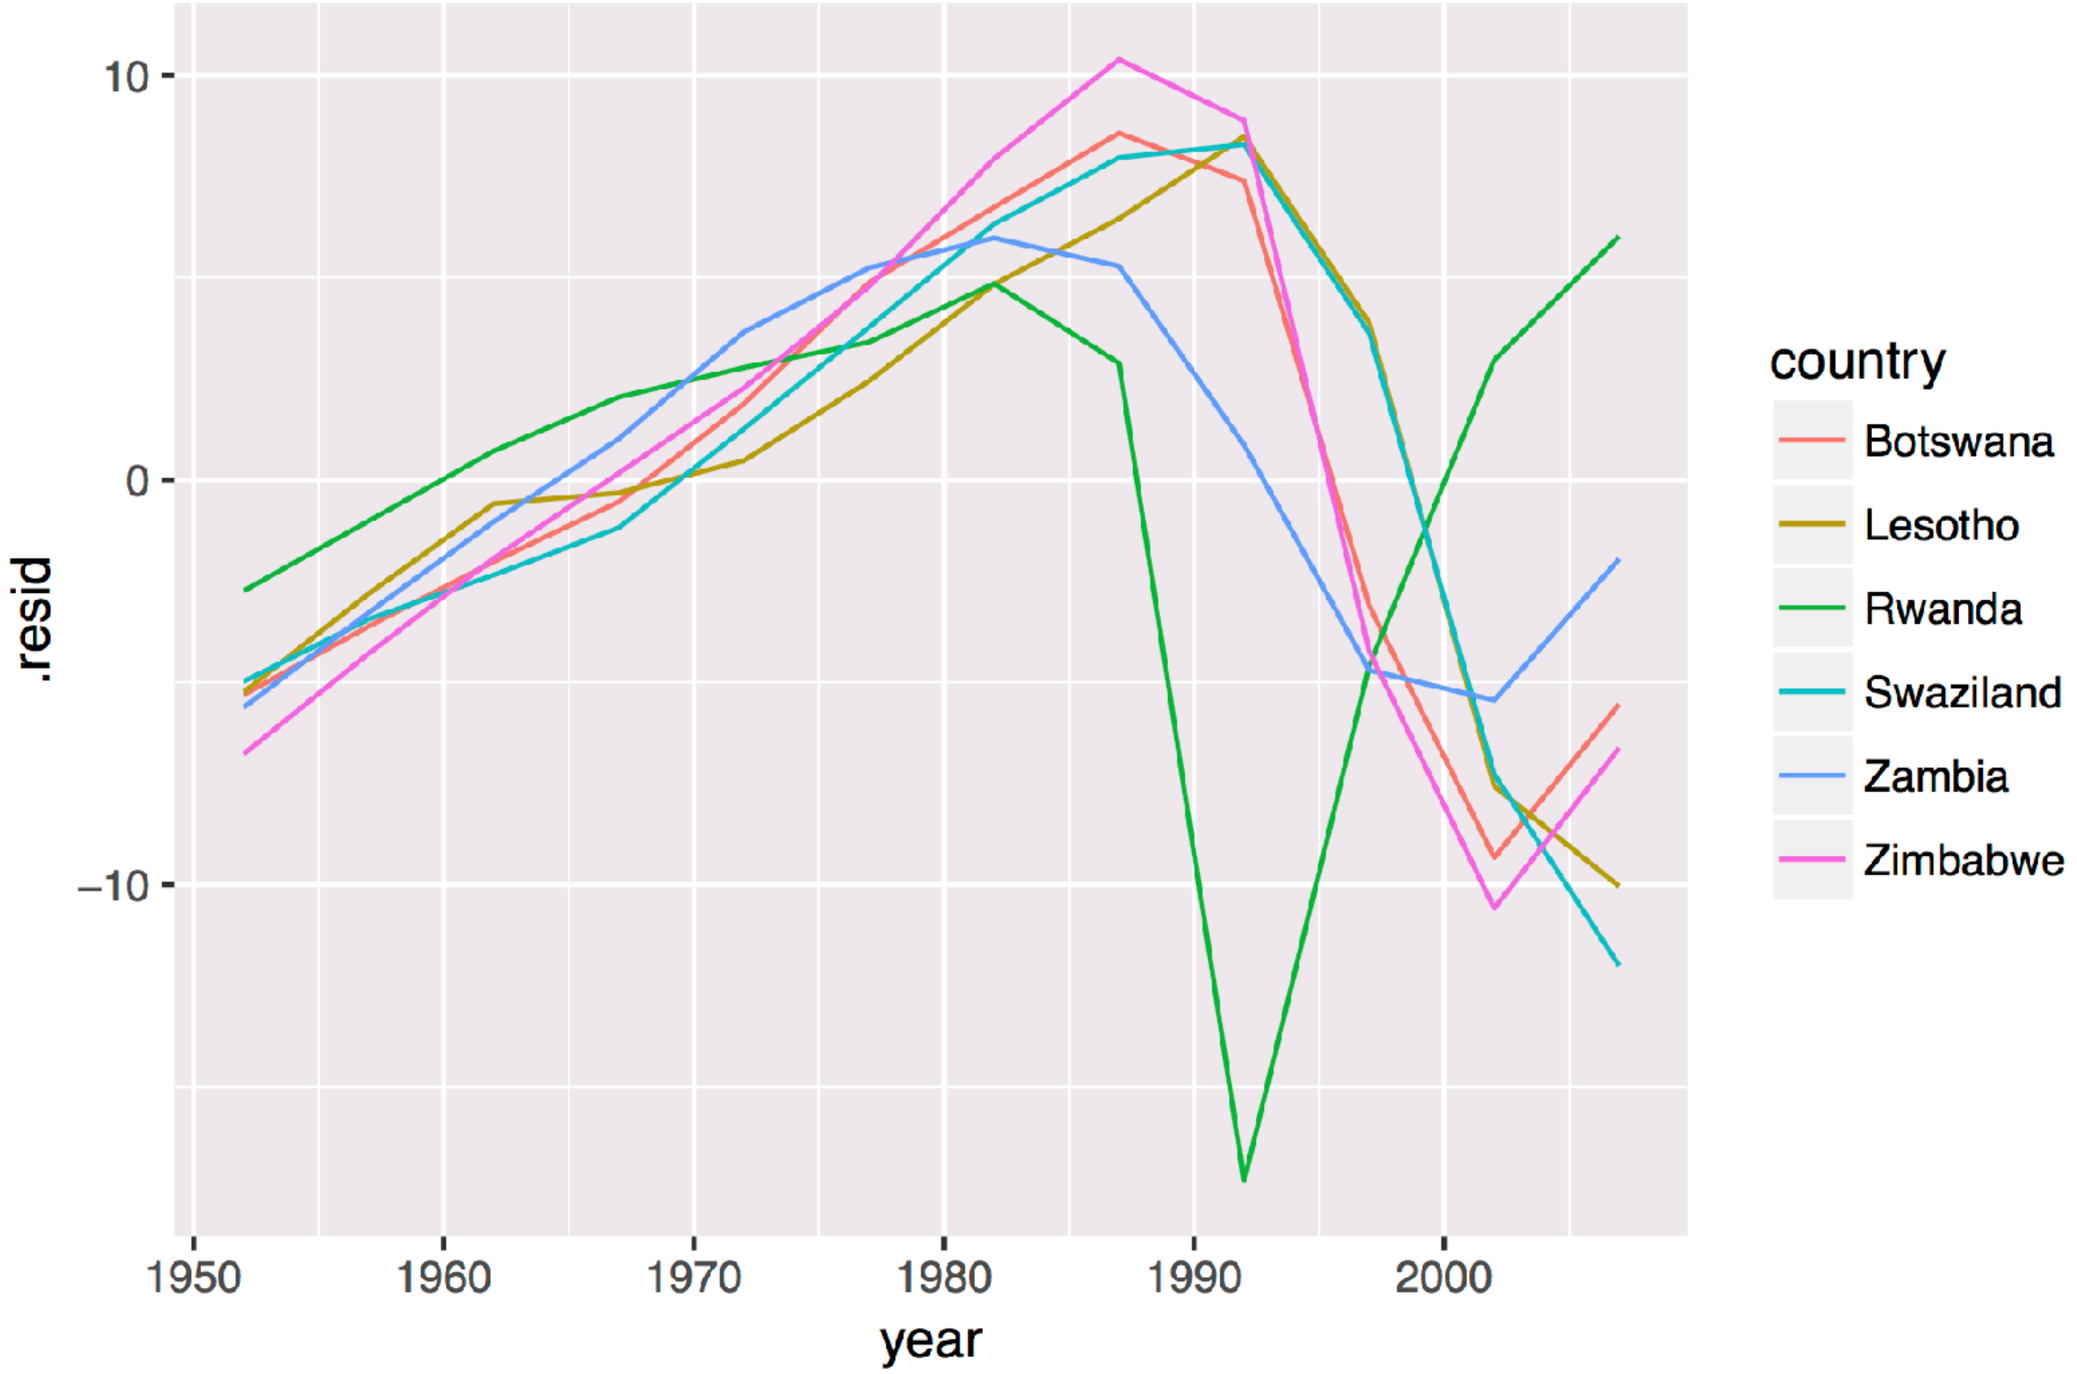

```
bad_fits <- gapminder %>%
  group_by(country, continent) %>%
  do(lm(lifeExp ~ year, data = .) %>%  glance()) %>%
  filter(r.squared < 0.25)

residuals <- gapminder %>%
  group_by(country, continent) %>%
  do(lm(lifeExp ~ year, data = .) %>%  augment())

residuals %>%
  semi_join(bad_fits) %>%
  ggplot() +
    geom_line(aes(year, .resid, color = country))
```

Fits the same model twice (to 142 countries each time)

```
bad_fits <- gapminder %>%
  group_by(country, continent) %>%
  do(lm(lifeExp ~ year, data = .) %>%  glance()) %>%
  filter(r.squared < 0.25)

residuals <- gapminder %>%
  group_by(country, continent) %>%
  do(lm(lifeExp ~ year, data = .) %>%  augment())

residuals %>%
  semi_join(bad_fits) %>%
  ggplot() +
    geom_line(aes(year, .resid, color = country))
```

Fits the same model twice (to 142 countries each time)

Need to keep track of two separate data frames

List columns

bad_fits

country	r.squared
Botswana	0.03
Lesotho	0.08
Rwanda	0.02
Swaziland	0.07
Zambia	0.06
Zimbabwe	0.06

residuals

country	year	.resid
Botswana	1952	-5.3071154
Botswana	1957	-3.6144580
Botswana	1962	-2.0158007
Botswana	1967	-0.5411434
Botswana	1972	1.8815140
Botswana	1977	4.8731713
Botswana	1982	6.7348287
Botswana	1987	8.5694860
Botswana	1992	7.3891434
Botswana	1997	-3.1031993
Botswana	2002	-9.3285420
Botswana	2007	-5.5378846
Lesotho	1952	-5.2410256
Lesotho	1957	-2.8098543
Lesotho	1962	-0.5876830
Lesotho	1967	-0.3205117
Lesotho	1972	0.4766597
Lesotho	1977	2.4398310
Lesotho	1982	4.8320023
Lesotho	1987	6.4561737
Lesotho	1992	8.4833450
Lesotho	1997	3.8785163
Lesotho	2002	-7.5643124
Lesotho	2007	-10.0431410
Rwanda	1952	-2.7419487
Rwanda	1957	-1.0127914
Rwanda	1962	0.7163660

master

country	r.squared	data
Botswana	0.03	year .resid 1952 -5.3071154 1957 -3.6144580 1962 -2.0158007 1967 -0.5411434 1972 1.8815140 1977 4.8731713 1982 6.7348287 1987 8.5694860 1992 7.3891434 1997 -3.1031993 2002 -9.3285420 2007 -5.5378846
Lesotho	0.08	year .resid 1952 -5.2410256 1957 -2.8098543 1962 -0.5876830 1967 -0.3205117 1972 0.4766597 1977 2.4398310 1982 4.8320023 1987 6.4561737 1992 8.4833450 1997 3.8785163 2002 -7.5643124 2007 -10.0431410
		year .resid 1952 -2.7419487 1957 -1.0127914

master

country	r.squared	data	model
Botswana	0.03	year .resid 1952 -5.3071154 1957 -3.6144580 1962 -2.0158007 1967 -0.5411434 1972 1.8815140 1977 4.8731713 1982 6.7348287 1987 8.5694860 1992 7.3891434 1997 -3.1031993 2002 -9.3285420 2007 -5.5378846	Call: lm(formula = lifeExp ~ year, data = .) Coefficients: (Intercept) year -65.49586 0.06067
Lesotho	0.08	year .resid 1952 -5.2410256 1957 -2.8098543 1962 -0.5876830 1967 -0.3205117 1972 0.4766597 1977 2.4398310 1982 4.8320023 1987 6.4561737 1992 8.4833450 1997 3.8785163 2002 -7.5643124 2007 -10.0431410	Call: lm(formula = lifeExp ~ year, data = .) Coefficients: (Intercept) year -139.16529 0.09557
		year .resid 1952 -2.7419487 1957 -1.0127914	

Quiz

How is a data frame/tibble similar to a list?

A data frame/tibble is a list!

data frame

num	cha	log
1	"one"	TRUE
2	"two"	FALSE
3	"three"	FALSE

=

List

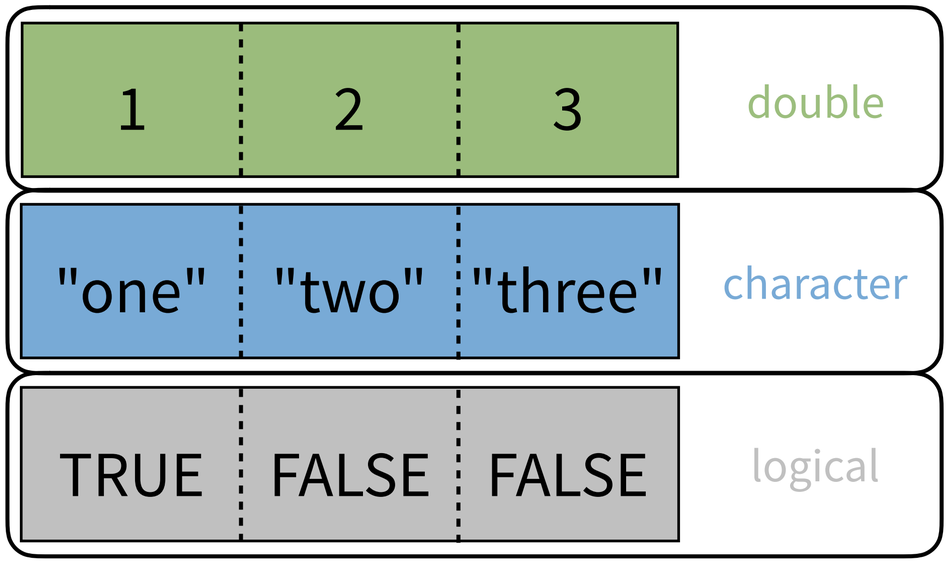

\+ class = "data.frame"

A data frame/tibble is a list!

data frame

num	cha	log
1	"one"	TRUE
2	"two"	FALSE
3	"three"	FALSE

`df["num"]`

num
1
2
3

`df[["num"]]`

`df$num`

`c(1, 2, 3)`

A data frame/tibble is a list!

data frame

num	cha	log
1	"one"	TRUE
2	"two"	FALSE
3	"three"	FALSE

`df %>% select(num)`

num
1
2
3

Quiz

If one of the elements of a list can be another list,

can one of the columns of a data frame be another list?

List

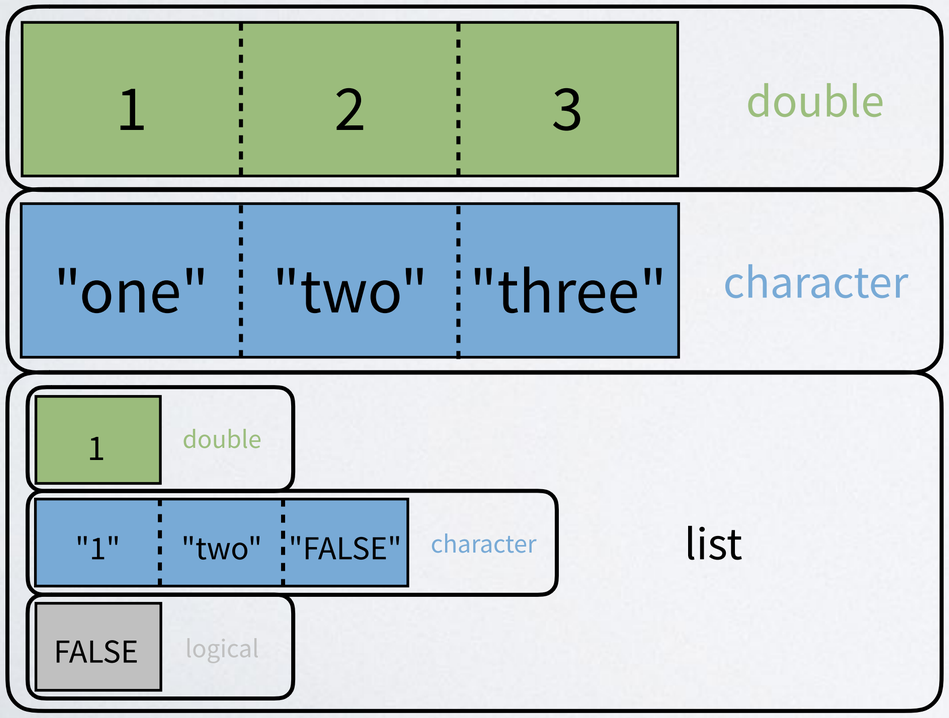

?

=

data frame

num	cha	listcol
1	"one"	1
2	"two"	c("1", "two", "FALSE")
3	"three"	FALSE

Yes.

```
tibble(
  num = c(1, 2, 3),
  cha = c("one", "two", "three"),
  listcol = list(1, c("1", "two", "FALSE"), FALSE)
)
```

num <dbl>	cha <chr>	listcol <list>
1	one	<dbl [1]>
2	two	<chr [3]>
3	three	<lgl [1]>

3 rows

nesting

nest()

Places grouped cases into a list column.

```
gapminder %>%
  group_by(country) %>%
  nest()
```

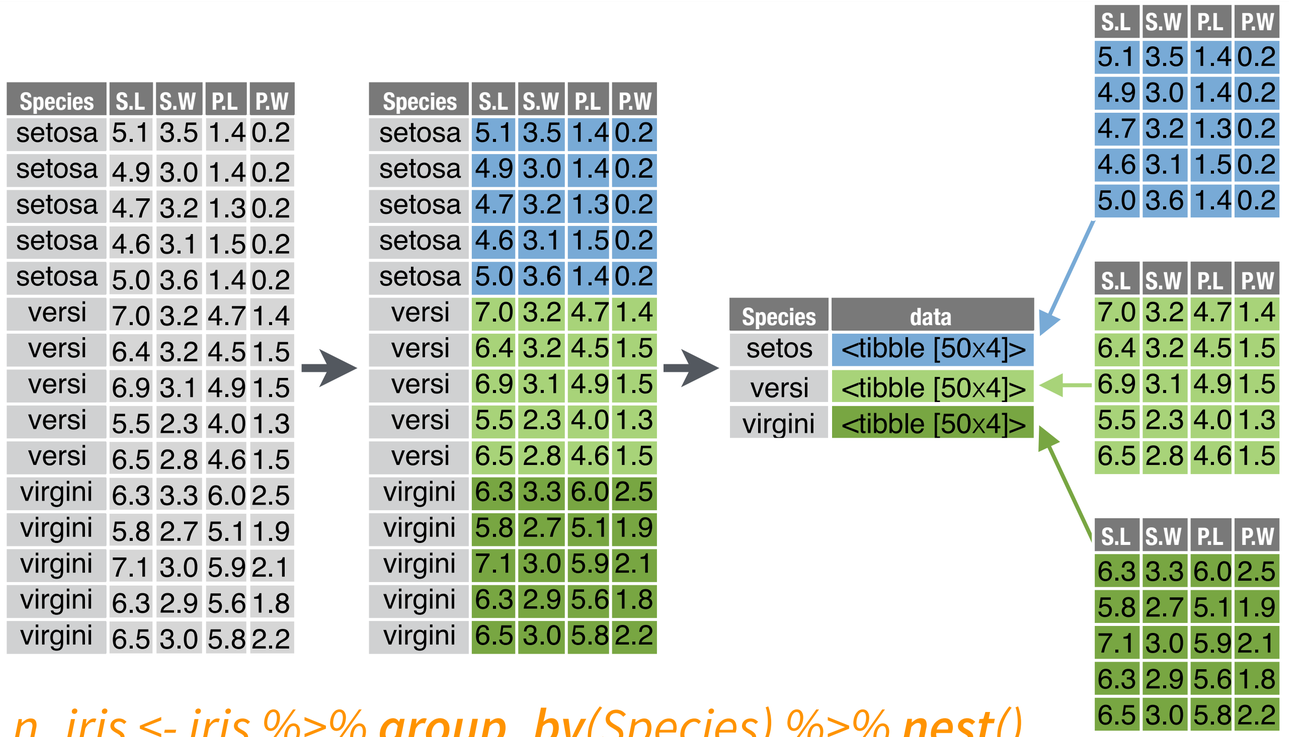

*n_iris <- iris %>% **group_by**(Species) %>% **nest**()*

nest()

```
residuals %>%
  group_by(country) %>%
  nest()
```

country	year	.resid
Botswana	1952	-5.3071154
Botswana	1957	-3.6144580
Botswana	1962	-2.0158007
Botswana	1967	-0.5411434
Botswana	1972	1.8815140
Botswana	1977	4.8731713
Botswana	1982	6.7348287
Botswana	1987	8.5694860
Botswana	1992	7.3891434
Botswana	1997	-3.1031993
Botswana	2002	-9.3285420
Botswana	2007	-5.5378846
Lesotho	1952	-5.2410256
Lesotho	1957	-2.8098543
Lesotho	1962	-0.5876830
Lesotho	1967	-0.3205117
Lesotho	1972	0.4766597
Lesotho	1977	2.4398310
Lesotho	1982	4.8320023
Lesotho	1987	6.4561737

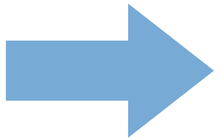

country	data
Botswana	year / .resid 1952 -5.3071154 1957 -3.6144580 1962 -2.0158007 1967 -0.5411434 1972 1.8815140 1977 4.8731713 1982 6.7348287 1987 8.5694860 1992 7.3891434 1997 -3.1031993 2002 -9.3285420 2007 -5.5378846
Lesotho	year / .resid 1952 -5.2410256 1957 -2.8098543 1962 -0.5876830 1967 -0.3205117

```
gapminder
```

country <fctr>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>
Afghanistan	Asia	1952	28.80100	8425333	779.4453
Afghanistan	Asia	1957	30.33200	9240934	820.8530
Afghanistan	Asia	1962	31.99700	10267083	853.1007
Afghanistan	Asia	1967	34.02000	11537966	836.1971
Afghanistan	Asia	1972	36.08800	13079460	739.9811
Afghanistan	Asia	1977	38.43800	14880372	786.1134
Afghanistan	Asia	1982	39.85400	12881816	978.0114
Afghanistan	Asia	1987	40.82200	13867957	852.3959
Afghanistan	Asia	1992	41.67400	16317921	649.3414
Afghanistan	Asia	1997	41.76300	22227415	635.3414

```
master <- gapminder %>%
  group_by(country) %>%
  nest()
```

country <fctr>	data <list>
Afghanistan	<tibble>
Albania	<tibble>
Algeria	<tibble>
Angola	<tibble>
Argentina	<tibble>
Australia	<tibble>
Austria	<tibble>
Bahrain	<tibble>
Bangladesh	<tibble>
Belgium	<tibble>

1–10 of 142 rows

Previous 1 2 3 4 5 6 ... 15 Next

```
master$data[[1]]
```

country <fctr>	data <list>
Afghanistan	<tibble>
Albania	<tibble>
Algeria	<tibble>
Angola	<tibble>
Argentina	<tibble>
Australia	<tibble>
Austria	<tibble>
Bahrain	<tibble>
Bangladesh	<tibble>
Belgium	<tibble>

1–10 of 142 rows

Previous 1 2 3 4 5 6 ... 15 Next

continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>
Asia	1952	28.801	8425333	779.4453
Asia	1957	30.332	9240934	820.8530
Asia	1962	31.997	10267083	853.1007
Asia	1967	34.020	11537966	836.1971
Asia	1972	36.088	13079460	739.9811
Asia	1977	38.438	14880372	786.1134
Asia	1982	39.854	12881816	978.0114
Asia	1987	40.822	13867957	852.3959
Asia	1992	41.674	16317921	649.3414
Asia	1997	41.763	22227415	635.3414

1–10 of 12 rows

Previous 1 2 Next

```
fit_model <- function(df) lm(lifeExp ~ year, data = df)

master <- master %>%
  mutate(model = map(data, fit_model))
```

country <fctr>	data <list>	model <list>
Afghanistan	<tibble>	<S3: lm>
Albania	<tibble>	<S3: lm>
Algeria	<tibble>	<S3: lm>
Angola	<tibble>	<S3: lm>
Argentina	<tibble>	<S3: lm>
Australia	<tibble>	<S3: lm>
Austria	<tibble>	<S3: lm>
Bahrain	<tibble>	<S3: lm>
Bangladesh	<tibble>	<S3: lm>
Belgium	<tibble>	<S3: lm>

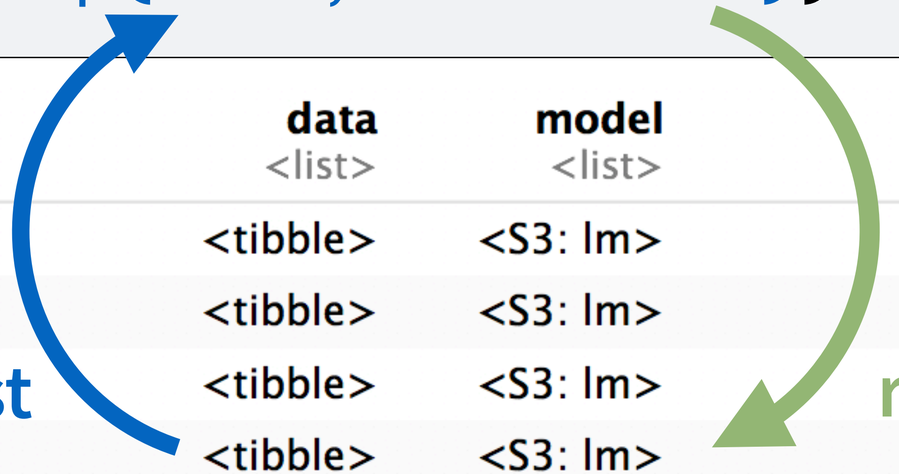

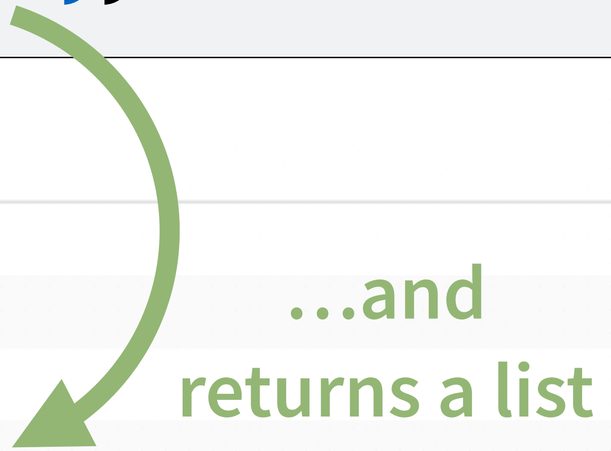

```
master$model[[1]]
```

country <fctr>	data <list>	model <list>
Afghanistan	<tibble>	<S3: lm>
Albania	<tibble>	<S3: lm>
Algeria	<tibble>	<S3: lm>
Angola	<tibble>	<S3: lm>
Argentina	<tibble>	<S3: lm>
Australia	<tibble>	<S3: lm>
Austria	<tibble>	<S3: lm>
Bahrain	<tibble>	<S3: lm>
Bangladesh	<tibble>	<S3: lm>
Belgium	<tibble>	<S3: lm>

```
Call:
lm(formula = lifeExp ~ year, data = x)

Coefficients:
(Intercept)         year
  -507.5343       0.2753
```

```
get_rsq <- function(mod) glance(mod)$r.squared

master <- master %>%
  mutate(r.squared = map_dbl(model, get_rsq))
```

country <fctr>	data <list>	model <list>	r.squared <dbl>
Afghanistan	<tibble>	<S3: lm>	0.94771226
Albania	<tibble>	<S3: lm>	0.91057777
Algeria	<tibble>	<S3: lm>	0.98511721
Angola	<tibble>	<S3: lm>	0.88781463
Argentina	<tibble>	<S3: lm>	0.99556810
Australia	<tibble>	<S3: lm>	0.97964774
Austria	<tibble>	<S3: lm>	0.99213401
Bahrain	<tibble>	<S3: lm>	0.96673981
Bangladesh	<tibble>	<S3: lm>	0.98936087
Belgium	<tibble>	<S3: lm>	0.99454056

map_dbl() takes a list

…and returns a number

Your Turn 4

Create your own copy of master and then add one more list column: **output** which contains the output of **augment()** for each model.

```
fit_model <- function(df) lm(lifeExp ~ year, data = df)
get_rsq <- function(mod) glance(mod)$r.squared
get_output <- function(mod) augment(mod)



master <- gapminder %>%
  group_by(country) %>%
  nest() %>%
  mutate(model = map(data, fit_model),
         r.squared = map_dbl(model, get_rsq),
         output = map(model, get_output))
master
```

country <fctr>	data <list>	model <list>	r.squared <dbl>	output <list>
Afghanistan	<tibble>	<S3: lm>	0.94771226	<data.frame [12 × 9]>
Albania	<tibble>	<S3: lm>	0.91057777	<data.frame [12 × 9]>
Algeria	<tibble>	<S3: lm>	0.98511721	<data.frame [12 × 9]>
Angola	<tibble>	<S3: lm>	0.88781463	<data.frame [12 × 9]>
Argentina	<tibble>	<S3: lm>	0.99556810	<data.frame [12 × 9]>
Australia	<tibble>	<S3: lm>	0.97964774	<data.frame [12 × 9]>
Austria	<tibble>	<S3: lm>	0.99213401	<data.frame [12 × 9]>
Bahrain	<tibble>	<S3: lm>	0.96673981	<data.frame [12 × 9]>
Bangladesh	<tibble>	<S3: lm>	0.98936087	<data.frame [12 × 9]>
Belgium	<tibble>	<S3: lm>	0.99454056	<data.frame [12 × 9]>

1–10 of 142 rows

```
master$output[[1]]
```

country <fctr>	data <list>	model <list>	r.squared <dbl>	output <list>
Afghanistan	<tibble>	<S3: lm>	0.94771226	<data.frame [12 × 9]>
Albania	<tibble>	<S3: lm>	0.91057777	<data.frame [12 × 9]>
Algeria				12 × 9]>
Angola				12 × 9]>
Argentina				12 × 9]>
Australia				12 × 9]>
Austria				12 × 9]>
Bahrain				12 × 9]>
Banglade				12 × 9]>
Belgium				12 × 9]>

1–10 of 1 … 15 Next

lifeExp <dbl>	year <int>	.fitted <dbl>	.se.fit <dbl>	.resid <dbl>	.hat <dbl>	.sigma <dbl>
28.801	1952	29.90729	0.6639995	-1.10629487	0.29487179	1.211813
30.332	1957	31.28394	0.5799442	-0.95193823	0.22494172	1.237512
31.997	1962	32.66058	0.5026799	-0.66358159	0.16899767	1.265886
34.020	1967	34.03722	0.4358337	-0.01722494	0.12703963	1.288917
36.088	1972	35.41387	0.3848726	0.67413170	0.09906760	1.267003
38.438	1977	36.79051	0.3566719	1.64748834	0.08508159	1.154002
39.854	1982	38.16716	0.3566719	1.68684499	0.08508159	1.147076
40.822	1987	39.54380	0.3848726	1.27820163	0.09906760	1.208243
41.674	1992	40.92044	0.4358337	0.75355828	0.12703963	1.260583
41.763	1997	42.29709	0.5026799	-0.53408508	0.16899767	1.274051

1–10 of 12 rows | 1–7 of 9 columns Previous 1 2 Next

Benefits

Data and models stay in correspondence across manipulations

```
master %>% filter(str_sub(country, 1, 1)  == "P")
```

country <fctr>	data <list>	model <list>	r.squared <dbl>	output <list>
Pakistan	<tibble>	<S3: lm>	0.9972497	<data.frame [12 × 9]>
Panama	<tibble>	<S3: lm>	0.9511952	<data.frame [12 × 9]>
Paraguay	<tibble>	<S3: lm>	0.9829865	<data.frame [12 × 9]>
Peru	<tibble>	<S3: lm>	0.9884740	<data.frame [12 × 9]>
Philippines	<tibble>	<S3: lm>	0.9914226	<data.frame [12 ×
Poland	<tibble>	<S3: lm>	0.8396631	<data.frame [1
Portugal	<tibble>	<S3: lm>	0.9690351	<data.fram 9]>
Puerto Rico	<tibble>	<S3: lm>	0.9078191	<data.f × 9]>

8 rows

But how can we plot?

unnest()

Unnests one or more list columns

```
master %>%
  unnest(data, output, .drop = FALSE)
```

A nested data frame

list columns to unnest (should contain data frames)

Drop remaining list columns from the result?

unnest()

Unnests one or more list columns

```
master %>%
  unnest(data, output, .drop = FALSE)
```

country <fctr>	r.squared <dbl>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>
Afghanistan	0.94771226	Asia	1952	28.80100	8425333	779.4453
Afghanistan	0.94771226	Asia	1957	30.33200	9240934	820.8530
Afghanistan	0.94771226	Asia	1962	31.99700	10267083	853.1007
Afghanistan	0.94771226	Asia	1967	34.02000	11537966	836.1971
Afghanistan	0.94771226	Asia	1972	36.08800	13079460	739.9811
Afghanistan	0.94771226	Asia	1977	38.43800	14880372	786.1134
Afghanistan	0.94771226	Asia	1982	39.85400	12881816	978.0114
Afghanistan	0.94771226	Asia	1987	40.82200	13867957	852.3959
Afghanistan	0.94771226	Asia	1992	41.67400	16317921	649.3414
Afghanistan	0.94771226	Asia	1997	41.76300	22227415	635.3414

1–10 of 1,704 rows | 1–7 of 16 columns Previous 1 2 3 4 5 6 ... 100 Next

```
master %>%
  filter(str_detect(country, pattern = "Korea")) %>%
  unnest(data)
```

country <fctr>	r.squared <dbl>	continent <fctr>	year <int>	lifeExp <dbl>	pop <int>	gdpPercap <dbl>
Korea, Dem. Rep.	0.7030631	Asia	1952	50.056	8865488	1088.278
Korea, Dem. Rep.	0.7030631	Asia	1957	54.081	9411381	1571.135
Korea, Dem. Rep.	0.7030631	Asia	1962	56.656	10917494	1621.694
Korea, Dem. Rep.	0.7030631	Asia	1967	59.942	12617009	2143.541
Korea, Dem. Rep.	0.7030631	Asia	1972	63.983	14781241	3701.622
Korea, Dem. Rep.	0.7030631	Asia	1977	67.159	16325320	4106.301
Korea, Dem. Rep.	0.7030631	Asia	1982	69.100	17647518	4106.525
Korea, Dem. Rep.	0.7030631	Asia	1987	70.647	19067554	4106.492
Korea, Dem. Rep.	0.7030631	Asia	1992	69.978	20711375	3726.064
Korea, Dem. Rep.	0.7030631	Asia	1997	67.727	21585105	1690.757

1–10 of 24 rows

Previous 1 2 3 Next

```
master %>%
  filter(str_detect(country, pattern = "Korea")) %>%
  unnest(data) %>%
  ggplot() +
    geom_line(aes(year, lifeExp, color = country))
```

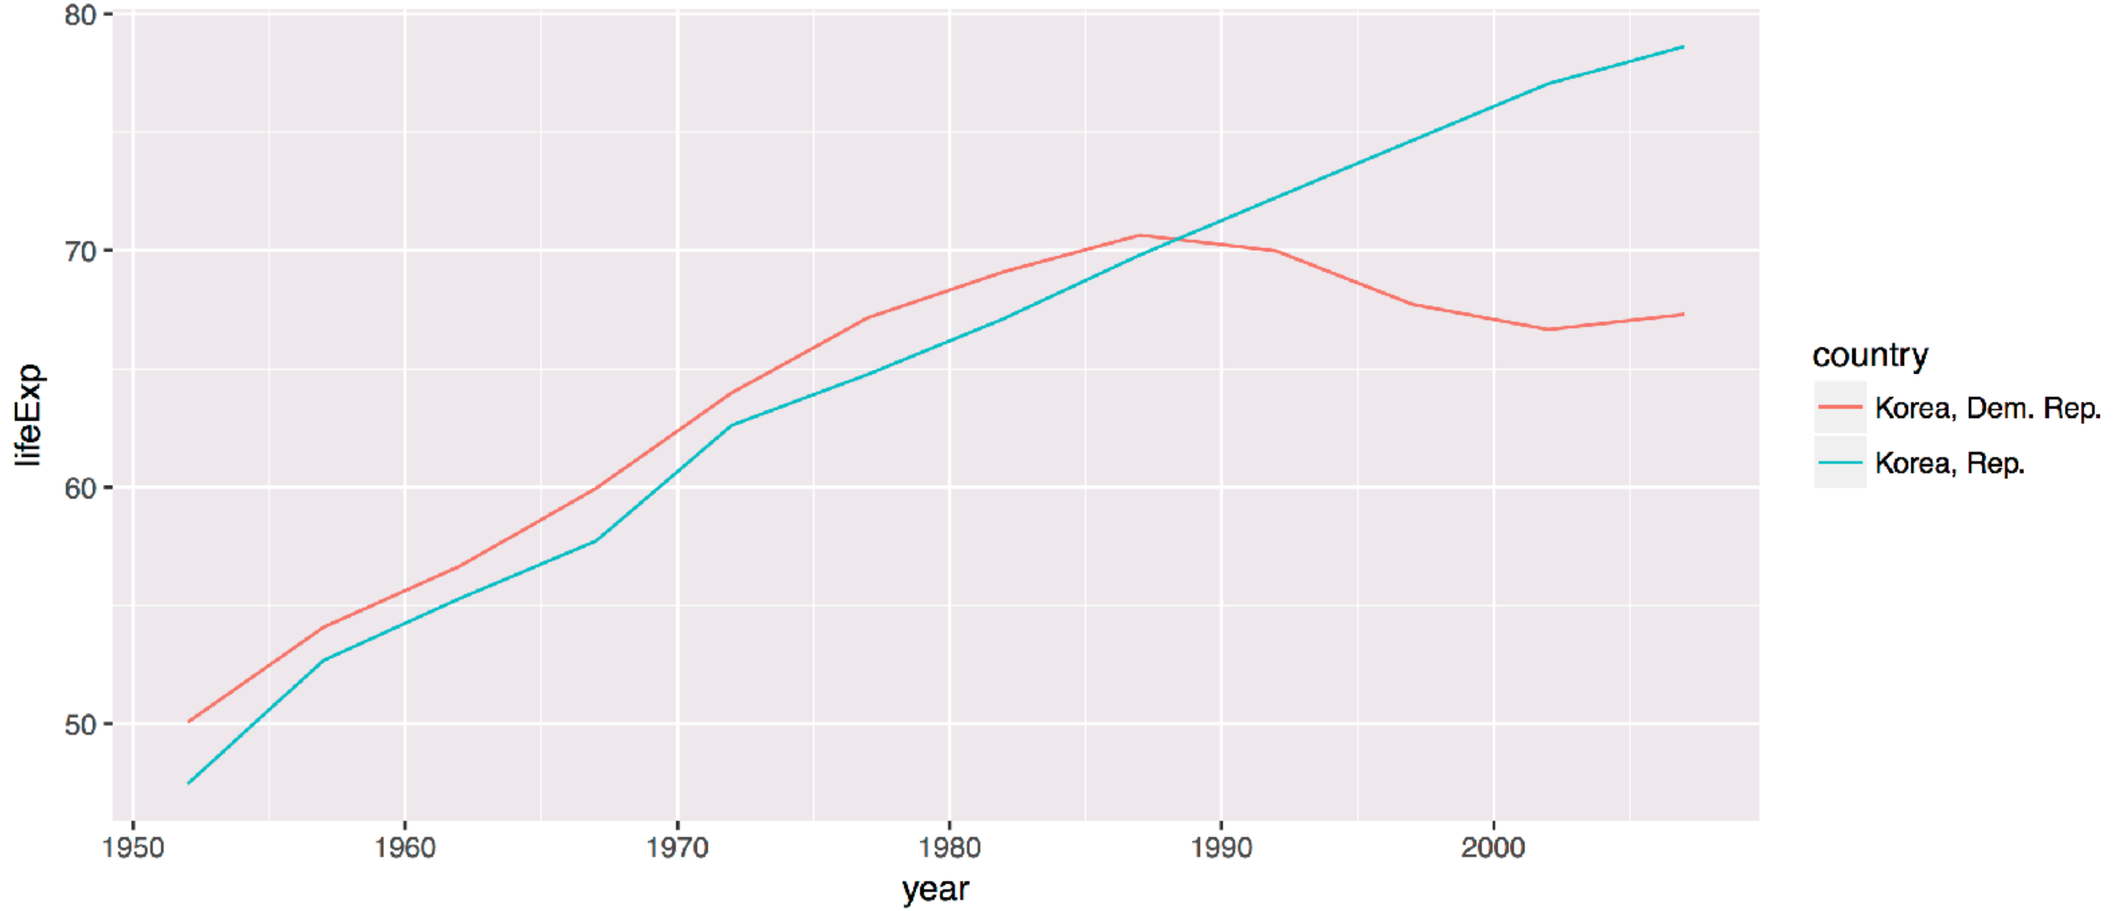

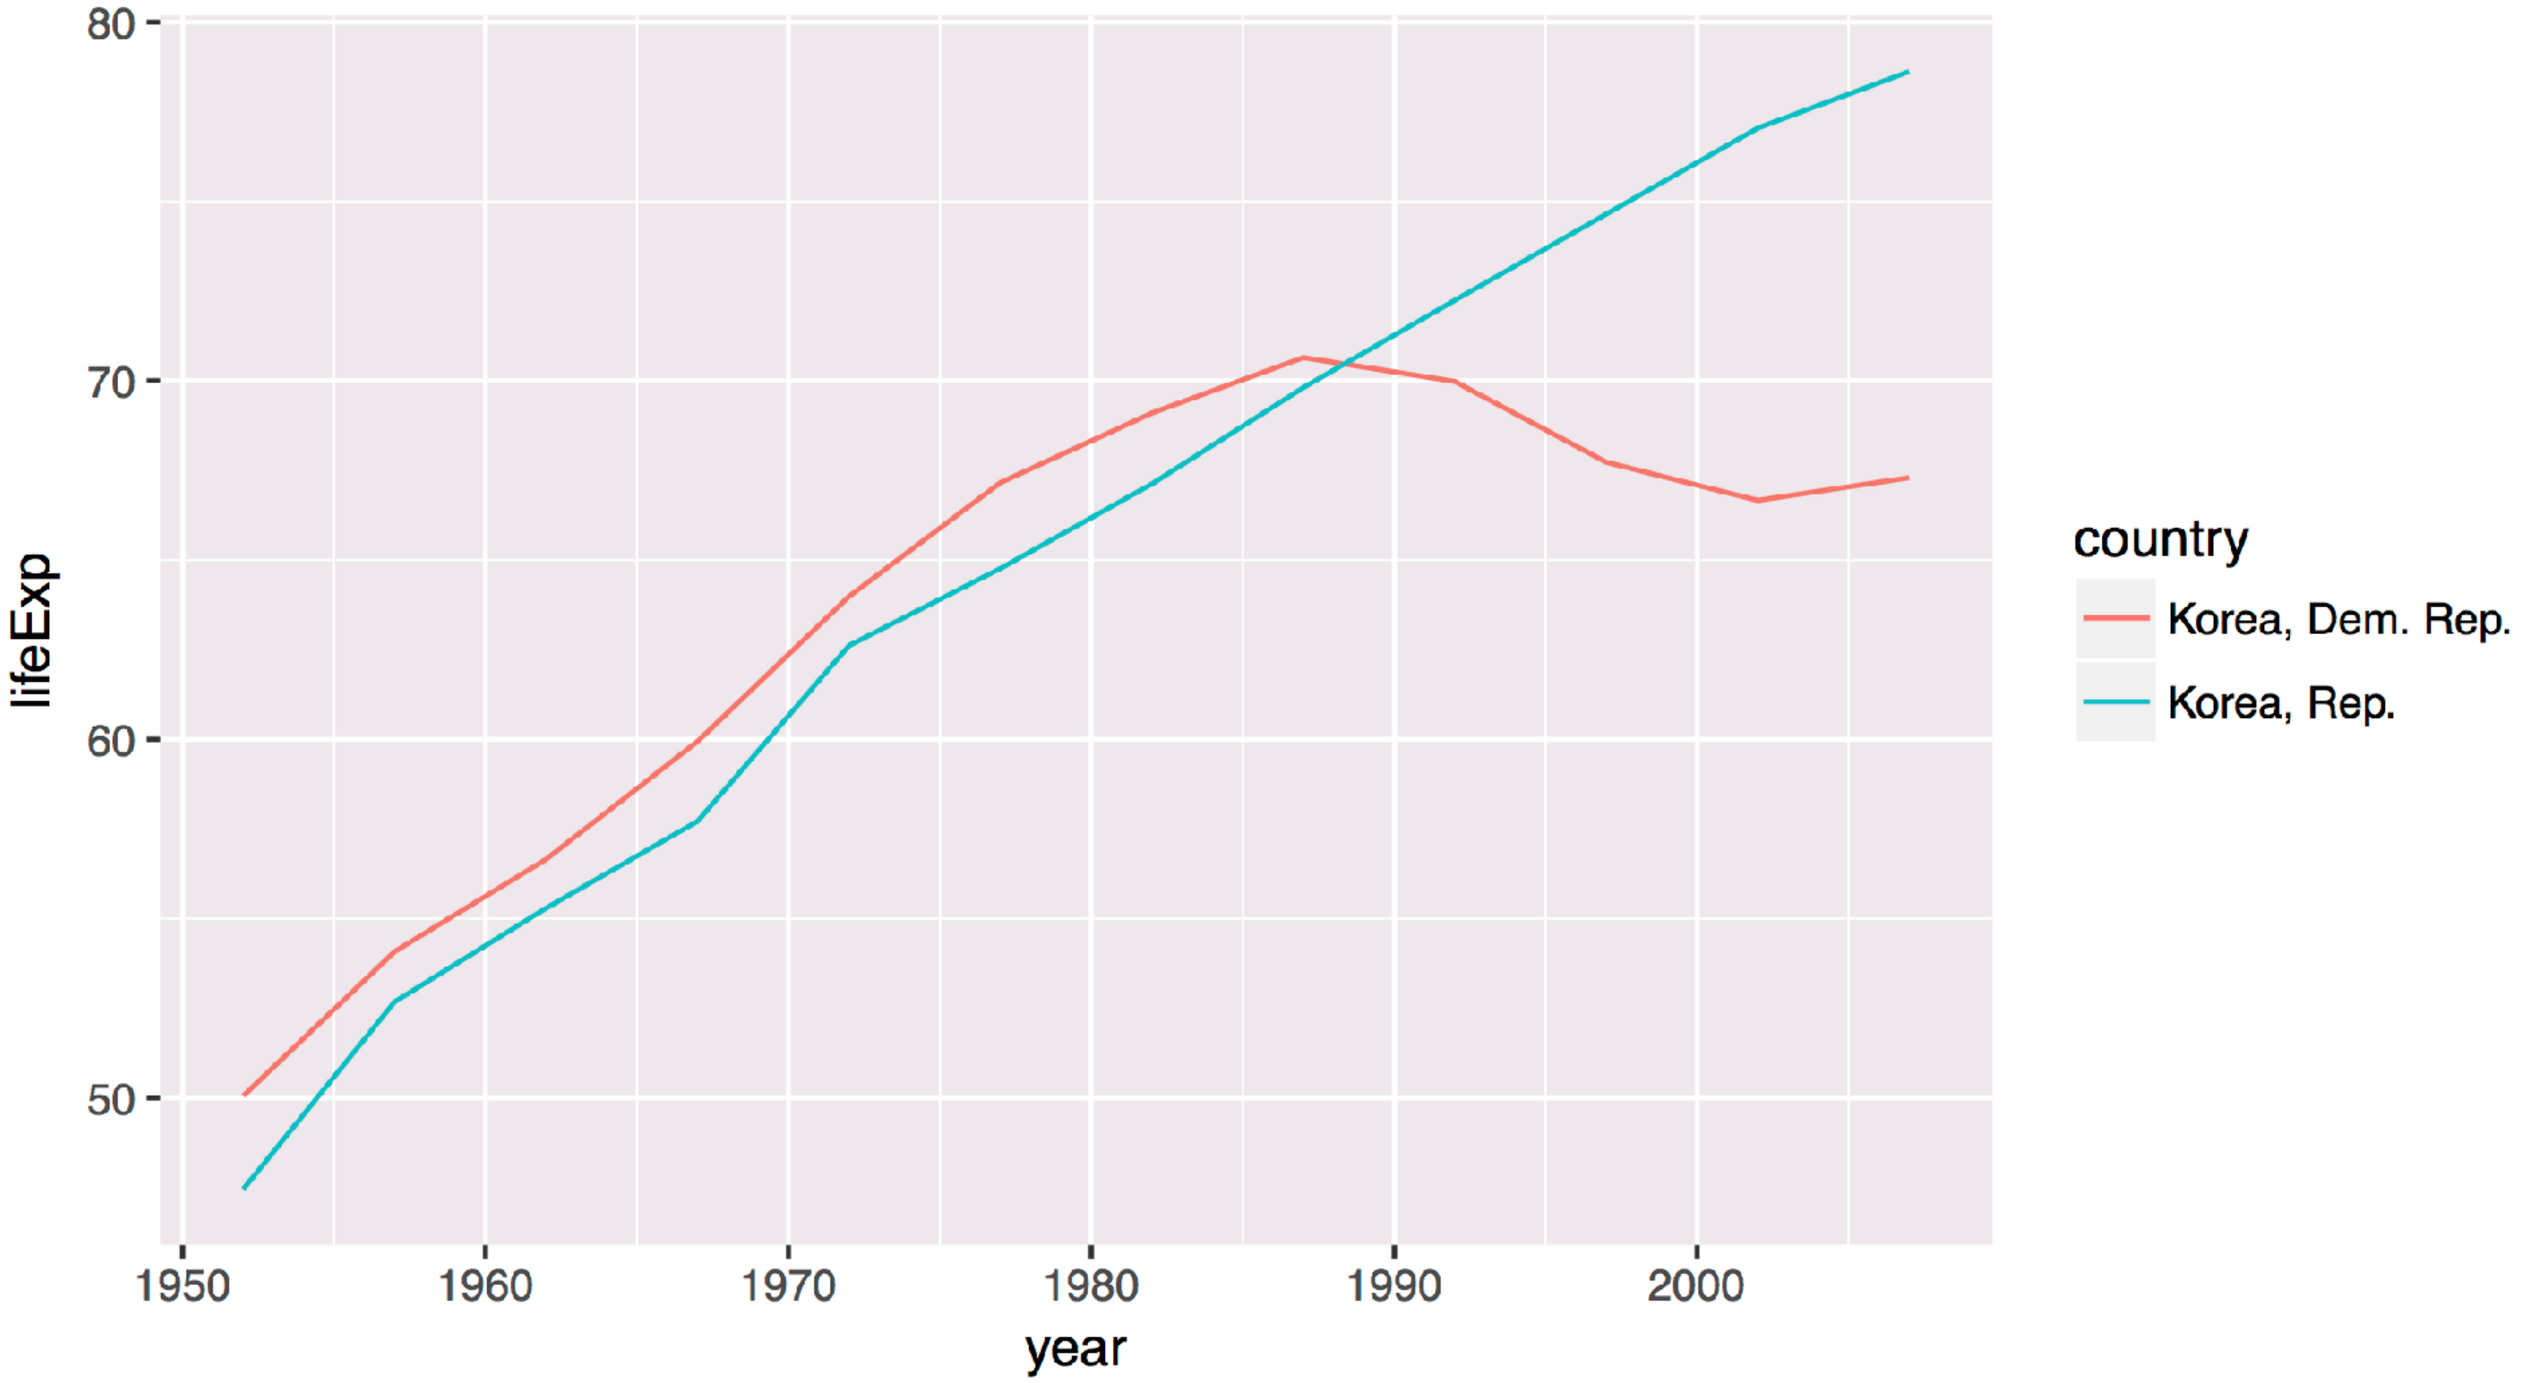
80
70
60
50
lifeExp
1950
1960
1970
1980
1990
2000
year
country
Korea, Dem. Rep.
Korea, Rep.

Your Turn 5

Use master to recreate our plot of the **residuals** vs **year** for the six countries with an r squared less than 0.25.

```
master %>%
  filter(r.squared < 0.25) %>%
  unnest(output) %>%
  ggplot() +
    geom_line(mapping = aes(x = year, y = .resid, color = country))
```

Take Away

A table is …an organizational structure …that you can manipulate.

country	r.squared	data	model
Botswana	0.03	year .resid 1952 -5.3071154 1957 -3.6144580 1962 -2.0158007 1967 -0.5411434 1972 1.8815140 1977 4.8731713 1982 6.7348287 1987 8.5694860 1992 7.3891434 1997 -3.1031993 2002 -9.3285420 2007 -5.5378846	Call: lm(formula = lifeExp ~ year, data = .) Coefficients: (Intercept) year -65.49586 0.06067
Lesotho	0.08	year .resid 1952 -5.2410256 1957 -2.8098543 1962 -0.5876830 1967 -0.3205117 1972 0.4766597 1977 2.4398310 1982 4.8320023 1987 6.4561737 1992 8.4833450 1997 3.8785163 2002 -7.5643124 2007 -10.0431410	Call: lm(formula = lifeExp ~ year, data = .) Coefficients: (Intercept) year -139.16529 0.09557

"Better experimental design = simpler statistics.
Better data model = simpler analysis."

- Jenny Bryan (2016)

Modeling with

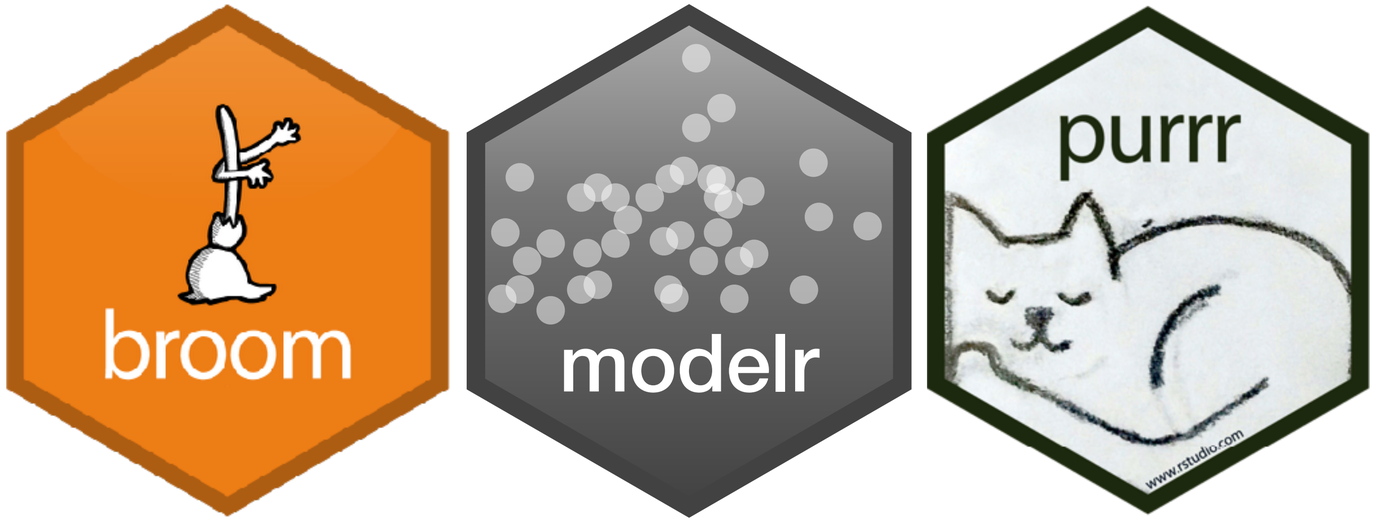

Thank You

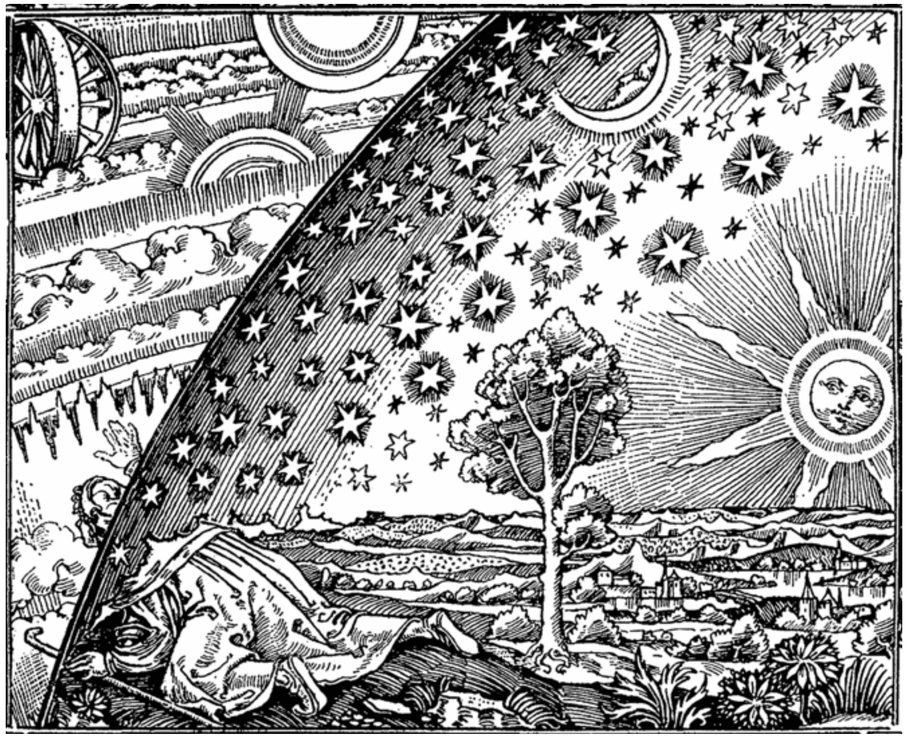

Please take the class survey

www.surveymonkey.com/r/SX9X69R